

Workpackage 5: High Performance Mathematical Computing

Clément Pernet

First OpenDreamKit Project review

Brussels, April 25, 2017

Goal: delivering high performance to maths users

Harnessing modern hardware \rightsquigarrow parallelisation

- ▶ in-core parallelism (SIMD vectorisation)
- ▶ multi-core parallelism
- ▶ distributed computing: clusters, cloud

Systems :

GAP

PARI/GP

SageMath

Singular

Components :

MPIR

LinBox

NumPy

Languages :

Cython

Python

Pythran

C

Architectures :

SIMD

**Multicore
server**

HPC cluster

Cloud

Goal: delivering high performance to maths users

Languages

- ▶ Computational Maths software use high level languages (e.g. Python)
 - ▶ High performance delivered by languages close to the metal (C, assembly)
- ⇒ compilation, automated optimisation

Systems :

GAP

PARI/GP

SageMath

Singular

Components :

MPIR

LinBox

NumPy

Languages :

Cython

Python

Pythran

C

Architectures :

SIMD

**Multicore
server**

HPC cluster

Cloud

Outline

Main tasks under review for the period

- Task 5.4: Singular
- Task 5.5: MPIR
- Task 5.6: Combinatorics
- Task 5.7: Pythran
- Task 5.8: SunGridEngine in JupyterHub

Progress report on other tasks

Singular: A computer algebra system for polynomial computations.

- ▶ Already has a generic parallelization framework
- ▶ Focus on optimizing kernel routines for fine grain parallelism

D5.6: Quadratic sieving for integer factorization

D5.7: Parallelization of matrix fast Fourier Transform

D5.6: Quadratic Sieving for integer factorization

Quadratic Sieving for integer factorization

Problem: Factor an integer n into prime factors

Role: Crucial in algebraic number theory, arithmetic geometry.

Earlier status: no HPC implementation for large instances:

- ▶ only fast code for up to 17 digits,
- ▶ only partial sequential implementation for large numbers

D5.6: Quadratic Sieving for integer factorization

Achievements

- ▶ Completed and debugged implementation of large prime variant
- ▶ Parallelised sieving component of implementation using OpenMP
- ▶ Experimented with a parallel implementation of Block Wiedemann

Results

- ▶ Now modern, robust, parallel code for numbers in 17–90 digit range

D5.6: Quadratic Sieving for integer factorization

Achievements

- ▶ Completed and debugged implementation of large prime variant
- ▶ Parallelised sieving component of implementation using OpenMP
- ▶ Experimented with a parallel implementation of Block Wiedemann

Results

- ▶ Now modern, robust, parallel code for numbers in 17–90 digit range
- ▶ Significantly faster on small multicore machines

Table: Speedup for four cores (c/f single core):

Digits	50	60	70	80	90
Speedup	1.1×	1.76×	1.55×	2.69×	2.80×

D5.7: Parallelise and assembly optimize FFT

FFT: Fast Fourier Transform over $\mathbb{Z}/p\mathbb{Z}$

- ▶ Among the top 10 most important algorithms
- ▶ Key to fast arithmetic (integers, polynomials)
- ▶ Difficult to optimize: high memory bandwidth requirement

Earlier status:

- ▶ world leading **sequential** code in MPIR and FLINT;
- ▶ no parallel code.

D5.7: Parallelise and assembly optimize FFT

Achievements

- ▶ Parallelised Matrix Fourier implementation using OpenMP
- ▶ Assembly optimised butterfly operations in MPIR

Results:

- ▶ $\approx 15\%$ speedup on Intel Haswell
- ▶ $\approx 20\%$ speedup on Intel Skylake
- ▶ Significant speedups on multicore machines

Table: Speedup of large integer multiplication on 4/8 cores:

Digits	3M	10M	35M	125M	700M	3.3B	14B
4 cores	1.35×	2.67×	2.92×	2.92×	3.01×	2.95×	3.32×
8 cores	1.35×	3.56×	4.22×	4.36×	4.50×	4.31×	5.49×

MPIR : a library for big integer arithmetic

- ▶ Bignum operations: fundamental across all of computer algebra

D5.5: Assembly superoptimization

- ▶ MPIR contains assembly language routines for bignum operations
 - ↪ hand optimised for every new microprocessor architecture
 - ↪ $\approx 3 - 6$ months of work for each architecture
- ▶ Superoptimisation: rearranges instructions to get optimal ordering

Earlier status:

- ▶ No assembly code for recent (> 2012) Intel and AMD chips (Bulldozer, Haswell, Skylake, ...)

D5.5: Assembly superoptimisation

Achievements

- ▶ A new assembly superoptimiser supporting recent instruction sets
- ▶ Superoptimised handwritten assembly code for Haswell and Skylake
- ▶ Hand picked faster assembly code for Bulldozer from existing implementations

Results:

- ▶ Sped up basic arithmetic operations for Bulldozer, Skylake and Haswell
- ▶ Noticeable speedups for bignum arithmetic for all size ranges

Op	Mul (s)	Mul (m)	Mul (b)	GCD (s)	GCD (m)	GCD (b)
Haswell	1.18×	1.27×	1.29×	0.72×	1.45×	1.27×
Skylake	1.15×	1.20×	1.22×	0.84×	1.65×	1.32×

s = 512 bits, m = 8192 bits, big = 100K bits

Task 5.6: Combinatorics

Perform a **map/reduce** operation on a very large set described **recursively**.

Large range of intensive applications in combinatorics:

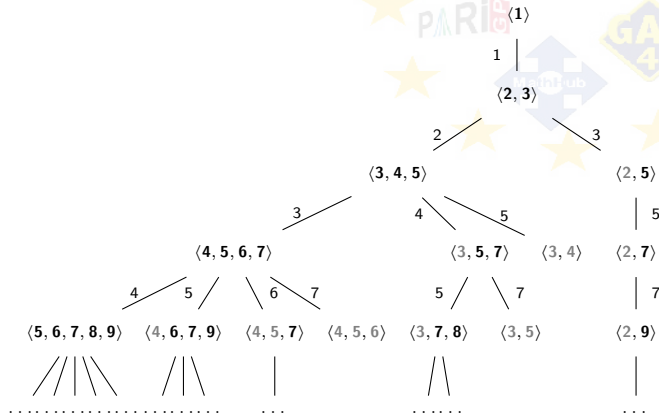
- ▶ Compute the cardinality, or more generally any kind of generating series
- ▶ Test a conjecture: i.e. find an element of S satisfying a specific property
- ▶ Count/list the elements of S having this property

Specificity of combinatorics:

- ▶ Typically the sets *don't fit in the computers memory* / disks and are enumerated on the fly (example of value: 10^{17} bytes).
- ▶ Easy to parallelize, if the set is flat (a list, a file, stored on a disk).

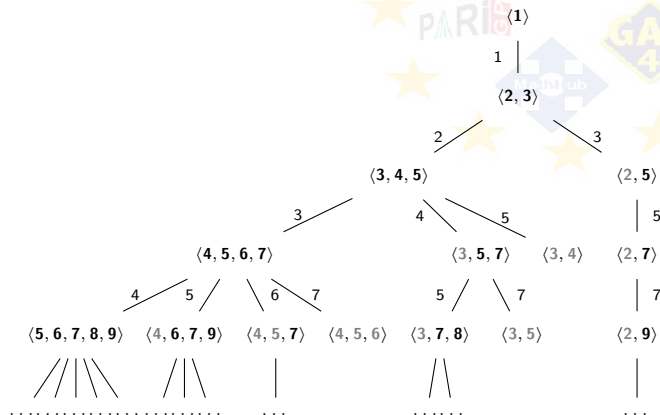
A Challenge: The tree of numerical semigroups

Extremely unbalanced. Need for an **efficient load balancing algorithm**.



A Challenge: The tree of numerical semigroups

Extremely unbalanced. Need for an **efficient load balancing algorithm**.



⇒ need for a high level task parallelization framework.

Work-Stealing System Architecture

A Python implementation

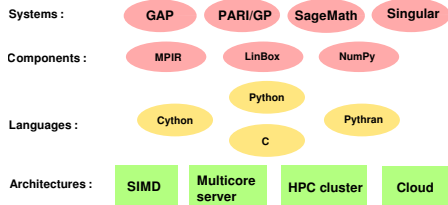
- ▶ Work stealing algorithm (Leiserson-Blumofe / Cilk)
- ▶ Easy to use, easy to call from sage
- ▶ Already, a dozen use case
- ▶ Scale well with the number of CPU cores
- ▶ Reasonably efficient (knowing that this is Python code).

References

- ▶ Trac Ticket 13580 <http://trac.sagemath.org/ticket/13580>
- ▶ *Exploring the Tree of Numerical Semigroups* Jean Fromentin and Florent Hivert <https://hal.inria.fr/UNIV-ROUEN/hal-00823339v3>

Task 5.7: Pythran

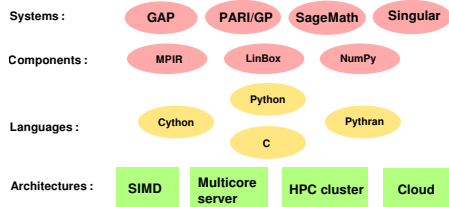
Pythran: a NumPy-centric Python to C compiler



- ▶ Many High level VRE rely on the Python language
- ▶ High performance is most often achieved by the C language

Task 5.7: Pythran

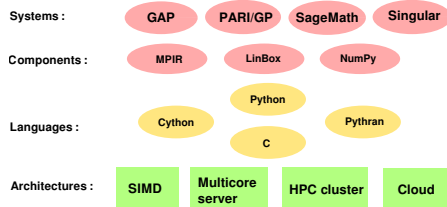
Pythran: a NumPy-centric Python to C compiler



- ▶ Many High level VRE rely on the Python language
- ▶ High performance is most often achieved by the C language
- ▶ Python to C compilers:
 - Cython:** general purpose
 - Pythran:** narrower scope, better at optimizing Numpy code (Linear algebra)

Task 5.7: Pythran

Pythran: a NumPy-centric Python to C compiler



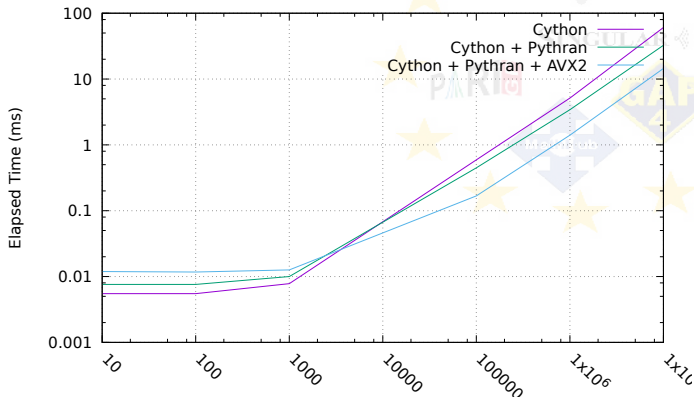
- ▶ Many High level VRE rely on the Python language
- ▶ High performance is most often achieved by the C language
- ▶ Python to C compilers:
 - Cython:** general purpose
 - Pythran:** narrower scope, better at optimizing Numpy code (Linear algebra)

Goal: Implement the convergence

D5.4 Improve Pythran typing system

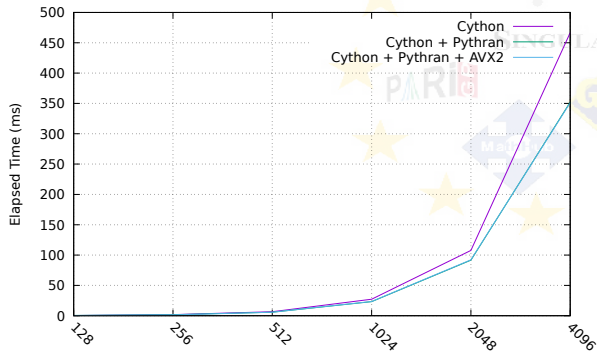
D5.2 Make Cython use Pythran backend to optimize Numpy code

D5.2: Make Cython use Pythran backend for NumPy code



```
import numpy
cimport numpy
def float_comp (numpy.ndarray[numpy.float_t, ndim=1] a,
                numpy.ndarray[numpy.float_t, ndim=1] b):
    return numpy.sqrt(numpy.sqrt(a*a+b*b))
```

D5.2: Make Cython use Pythran backend for NumPy code



```
def harris(numpy.ndarray[numpy.float_t, ndim=2] I):  
    cdef int m = I.shape[0]  
    cdef int n = I.shape[1]  
    cdef numpy.ndarray[numpy.float_t, ndim=2] dx = (I[1:, :] - I[:m-1, :])[1:, :]  
    cdef numpy.ndarray[numpy.float_t, ndim=2] dy = (I[:, 1:] - I[:, :n-1])[1:, :]  
    cdef numpy.ndarray[numpy.float_t, ndim=2] A = dx * dx  
    cdef numpy.ndarray[numpy.float_t, ndim=2] B = dy * dy  
    cdef numpy.ndarray[numpy.float_t, ndim=2] C = dx * dy  
    cdef numpy.ndarray[numpy.float_t, ndim=2] tr = A + B  
    cdef numpy.ndarray[numpy.float_t, ndim=2] det = A * B - C * C  
    return det - tr * tr
```

Task 5.8: SunGridEngine integration in JupyterHub

Access to big compute

- ▶ Traditional access to supercomputers is difficult
- ▶ Notebooks are easy but run on laptops or desktops
- ▶ We need a way to connect notebooks to supercomputers

Sun Grid Engine

A job scheduler for Academic HPC Clusters

- ▶ Controls how resources are allocated to researchers
- ▶ One of the most popular schedulers

Achievements: D5.3

- ▶ Developed software to run Jupyter notebooks on supercomputers
- ▶ Users don't need to know details. They just log in.
- ▶ Demonstration install at University of Sheffield

Outline

Main tasks under review for the period

- Task 5.4: Singular
- Task 5.5: MPIR
- Task 5.6: Combinatorics
- Task 5.7: Pythran
- Task 5.8: SunGridEngine in JupyterHub

Progress report on other tasks

Progress report on other tasks

T5.1: Pari

- ▶ Generic parallelization engine is now mature, released (D5.10, due M24)

T5.2: GAP

- ▶ 6 releases were cut integrating contributions of D3.11 and D5.15
- ▶ Build system refactoring for integration of HPC GAP

T5.3: LinBox

- ▶ Algorithmic advances (5 articles) on linear algebra and verified computing
- ▶ Software releases and integration into SageMath