

# Workpackage 5: High Performance Mathematical Computing

Vincent Delecroix,  
Alexander Konovalov,  
Clément Pernet,

Second OpenDreamKit Project review

Luxembourg, October 30, 2018

## Mathematical computing

Computing with a large variety of objects

►  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$
$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

for applications where all digits matter.

## Mathematical computing

Computing with a large variety of objects

- ▶  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over  $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

- ▶ Tree algebra

for applications where all digits matter.

$$\frac{3q^2 - q^5}{q^5 + 2q^4 + 3q^3 + 3q^2 + 2q + 1} \begin{array}{c} (a) \\ / \quad \backslash \\ (b) \quad (c) \\ \quad \quad | \\ \quad \quad (d) \end{array} + \frac{2q}{q^4 + q^3 + 2q^2 + q + 1} \begin{array}{c} (a) \\ | \\ (b) \\ / \quad \backslash \\ (c) \quad (d) \end{array}$$

# High performance mathematical computing

**Need for High performance:** applications where size matters:

**Experimental maths:** testing conjectures

- ▶ larger instances give higher confidence

# High performance mathematical computing

**Need for High performance:** applications where size matters:

**Experimental maths:** testing conjectures

- ▶ larger instances give higher confidence

**Algebraic cryptanalysis:** security = computational difficulty

- ▶ key size determined by the largest solvable problem

## Example

Breaking RSA by integer factorization:  $n = pq$ . Last record:

- ▶  $n$  of 768 bits
- ▶ linear algebra in dimension 192 796 550 over  $\mathbb{F}_2$  (105Gb)
- ▶ About 2000 CPU years



# High performance mathematical computing

**Need for High performance:** applications where size matters:

**Experimental maths:** testing conjectures

- ▶ larger instances give higher confidence

**Algebraic cryptanalysis:** security = computational difficulty

- ▶ key size determined by the largest solvable problem

## Example

Breaking RSA by integer factorization:  $n = pq$ . Last record:

- ▶  $n$  of 768 bits
- ▶ linear algebra in dimension 192 796 550 over  $\mathbb{F}_2$  (105Gb)
- ▶ About 2000 CPU years

**3D data analysis, shape recognition:**

- ▶ via persistent homology
- ▶ large sparse matrices over  $\mathbb{F}_2$ ,  $\mathbb{Z}$

# Goal: delivering high performance to maths users

**Systems :**

**GAP**

**PARI/GP**

**SageMath**

**Singular**

**Components :**

**FLINT**

**MPIR**

**LinBox**

**PPL**

**NumPy**

# Goal: delivering high performance to maths users

## Harnessing modern hardware $\rightsquigarrow$ parallelisation

- ▶ in-core parallelism (SIMD vectorisation)
- ▶ multi-core parallelism
- ▶ distributed computing: clusters, cloud

**Systems :**

**GAP**

**PARI/GP**

**SageMath**

**Singular**

**Components :**

**FLINT**

**MPIR**

**LinBox**

**PPL**

**NumPy**

**Architectures :**

**SIMD**

**Multicore  
server**

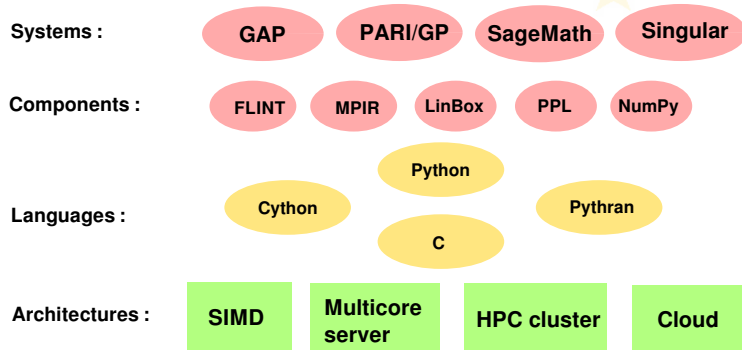
**HPC cluster**

**Cloud**

# Goal: delivering high performance to maths users

## Languages

- ▶ Computational Maths software uses high level languages (e.g. Python)
  - ▶ High performance delivered by languages close to the metal (C, assembly)
- ~> compilation, automated optimisation



# High performance mathematical computing

## Goal:

- ▶ Improve/Develop parallelization of software components
- ▶ Expose them through the software stack
- ▶ Offer High Performance Computing to VRE's users

## Milestone M8: Seamless use of parallel computing architecture in the VRE (proof of concept)

*Astrid wants to run compute intensive routines involving both dense linear algebra and combinatorics. She has access through a JupyterHub-based VRE to a high end multi-core machine which includes a vanilla SAGE installation. She automatically benefits from the HPC features of the underlying specialized libraries (LinBox, ...). This is a proof of concept of the overall framework to integrate the HPC advances of specialized libraries into a general purpose VRE. It will prepare the final integration of a broader set of such parallel features for the end of the project*

# Outline

## Workpackage management

### Exact linear algebra

Exact linear algebra algorithms and implementations (D5.12).

Distributed computing

### Combinatorics

Refactor and optimize Sage's Combinatorics (D5.11)

### Progress report on other tasks

T5.1: Pari

T5.2: GAP

T5.4: Singular

# Outline

## Workpackage management

### Exact linear algebra

Exact linear algebra algorithms and implementations (D5.12).

Distributed computing

### Combinatorics

Refactor and optimize Sage's Combinatorics (D5.11)

### Progress report on other tasks

T5.1: Pari

T5.2: GAP

T5.4: Singular

# Outcome of WorkPackage 5

Component	Review 1	Review 2	Final review
T5.1 Pari/GP			D5.16
T5.2 GAP			D5.15
T5.3 LinBox		D5.12	D5.14
T5.4 Singular	D5.6, D5.7		D5.13
T5.5 MPIR	D5.5, D5.7		
T5.6 Combinatorics	D5.1	D5.11	
T5.7 Pythran	D5.2	D5.11	
T5.8 SunGrid Engine	D5.3		

## Overall

- ▶ over 20 software releases were cut
- ▶ 7 research papers


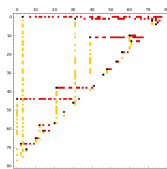
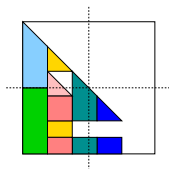


**Recommendation 10:** *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be nearer the leading edge.*

**Recommendation 10:** *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be **nearer the leading edge**.*

## Leading edge achievements in linear algebra

- ▶ symmetric factorization outperforms LAPACK implementation
- ▶ new non-hierarchical generator for quasiseparable matrices
- ▶ large scale parallelization of rational linear solver



# Strengthening interactions with numerical HPC community

## Existing connection

Dense linear algebra numerical BLAS used for exact FFLAS for 17 years

Pointwise interactions: J. Dongarra, L. Grigori, J-Y. L'Excellent, etc

Publications to in main HPC venues: SIAM-PPSC, EuroPar, PMAA, ParCo

Animation: involved in the French CNRS *Calcul* working group (Sci. Comp.)

# Strengthening interactions with numerical HPC community

## Existing connection

Dense linear algebra numerical BLAS used for exact FFLAS for 17 years

Pointwise interactions: J. Dongarra, L. Grigori, J-Y. L'Excellent, etc

Publications to in main HPC venues: SIAM-PPSC, EuroPar, PMAA, ParCo

Animation: involved in the French CNRS *Calcul* working group (Sci. Comp.)

## Recently established

- ▶ With the BLIS group:
  - ▶ Reporting bugs (SIMD vectorization)
  - ▶ Sharing experience in implementing Strassen's algorithm
- ▶ On-going collaboration with T. Mary (Mumps) and S. Chandrasekaran (UCSB) on quasiseparable matrix algorithmic

# Outline

Workpackage management

Exact linear algebra

Exact linear algebra algorithms and implementations (D5.12).

Distributed computing

Combinatorics

Refactor and optimize Sage's Combinatorics (D5.11)

Progress report on other tasks

T5.1: Pari

T5.2: GAP

T5.4: Singular

## Task 5.3: LinBox, High performance exact linear algebra

*Mathematics is the art of reducing any problem to linear algebra*

SINGULAR – W. Stein

PARIGe



## Task 5.3: LinBox, High performance exact linear algebra

*Mathematics is the art of reducing any problem to linear algebra*

SINGULAR – W. Stein

**Linear algebra: a key building block for HPC**

### Similarities with numerical HPC

- ▶ central elementary problem to which others reduce to
- ▶ (rather) simple algorithmic
- ▶ high compute/memory intensity

## Task 5.3: LinBox, High performance exact linear algebra

*Mathematics is the art of reducing any problem to linear algebra*

SINGULAR – W. Stein

**Linear algebra: a key building block for HPC**

### Similarities with numerical HPC

- ▶ central elementary problem to which others reduce to
- ▶ (rather) simple algorithmic
- ▶ high compute/memory intensity

### Specificities

- ▶ Multiprecision arithmetic  $\Rightarrow$  lifting from finite precision ( $\mathbb{F}_p$ )
- ▶ Rank deficiency  $\Rightarrow$  unbalanced dynamic blocking
- ▶ Early adopter of subcubic matrix arithmetic  $\Rightarrow$  recursion



## D5.12: Exact linear algebra algorithms and implementations. Library maintenance and close integration in mathematical software for LinBox library



SINGULAR

PARIG



LMFDB

### 1. Algorithmic innovations:

- 1.1 Rank deficient dense Gaussian elimination
- 1.2 Quasiseparable matrices
- 1.3 Outsourced computing security

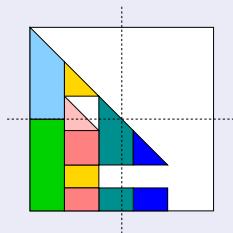
### 2. Software releases and integration:

- 2.1 LinBox ecosystem: LinBox, fflas-ffpack, givaro
- 2.2 SageMath integration



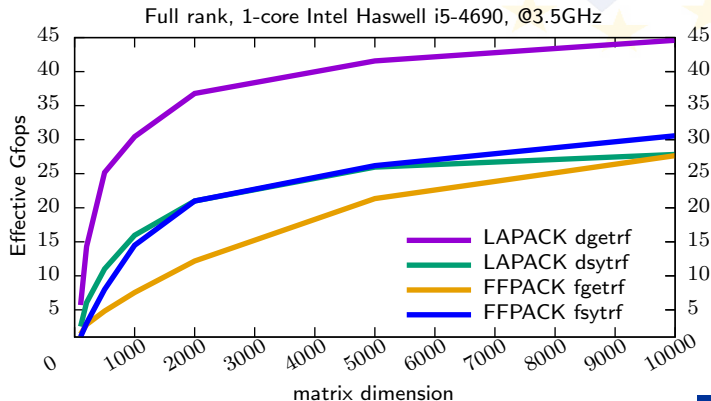
## [ISSAC'18] Symmetric triangular factorization

- ▶ First unconditional recursive algorithm
- ▶ Pivoting revealing the Rank Profile Matrix
- ▶  $O(n^2 r^{\omega-2})$  ( $= 1/3 n^3$  with  $\omega = 3, r = n$ )
- ▶ Also hot topic in numerical linalg  
(LAPACK Working notes 294, Dec'17)



# LAPACK vs FFPACK modulo 8 388 593

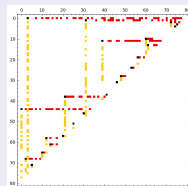
$n$	LAPACK		FFPACK	
	dgetrf (LU)	dsytrf (LDLT)	fgetrf (LU)	fsytrf (LDLT)
5000	2.01s	1.60s	3.90s	1.59s
10000	14.95s	11.98s	24.12s	10.90s



*Matrices with low off-diagonal rank*

[ISSAC'16, JSC'18] New compact representation and algorithms

- ▶ Matches the best space complexities
- ▶ Reduction to matrix multiplication
- ▶ **Breakthrough:** Flat representation (non hierarchical)



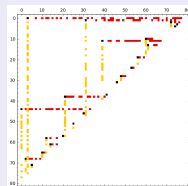
*Matrices with low off-diagonal rank*

[ISSAC'16, JSC'18] New compact representation and algorithms

- ▶ Matches the best space complexities
- ▶ Reduction to matrix multiplication
- ▶ **Breakthrough:** Flat representation (non hierarchical)

**Follow-up:** on-going collaboration with numerical HPC experts:

- ▶ S. Chandrasekaran (UCSB)
- ▶ T. Mary (U. Manchester, Mumps)



## LinBox ecosystem

`givaro`: field/ring arithmetic

`fflas-ffpack`: dense linear algebra over finite field

`LinBox`: exact linear algebra

Tightly integrated in SageMath

# Software releases and integration

## LinBox ecosystem

<code>givaro</code> : field/ring arithmetic	4 releases
<code>fflas-ffpack</code> : dense linear algebra over finite field	6 releases
<code>LinBox</code> : exact linear algebra	6 releases
Tightly integrated in SageMath	13 tickets

# Software releases and integration

## LinBox ecosystem

<code>givaro</code> : field/ring arithmetic	4 releases
<code>fflas-ffpack</code> : dense linear algebra over finite field	6 releases
<code>LinBox</code> : exact linear algebra	6 releases
Tightly integrated in SageMath	13 tickets

## Featuring

- ▶ Full functional implementations of new algorithmic contributions
- ▶ Improved vectorization and parallel routines
- ▶ Drastic improvement of reliability (continuous integration, test-suite coverage, randomized certificates, etc)



# Distributed computing (on-going work)



## D5.14: Distributed exact linear system solving

- ▶ 2 full time engineers
- ▶ Communication and serialization layer done
- ▶ Prototype MPI parallelization of Chinese remainder based solver.

# Distributed computing (on-going work)

## D5.14: Distributed exact linear system solving

- ▶ 2 full time engineers
- ▶ Communication and serialization layer done
- ▶ Prototype MPI parallelization of Chinese remainder based solver.





# Distributed computing (on-going work)

## D5.14: Distributed exact linear system solving

### Roadmap:

- ▶ Major refactorization of LinBox solver code under way
- ▶ Parallelization of Dixon-lifting solver
- ▶ Hybrid combination of CRT+Dixon.
- ▶ Hybrid OpenMP-MPI implementation



# Outline

Workpackage management

Exact linear algebra

Exact linear algebra algorithms and implementations (D5.12).

Distributed computing

Combinatorics

Refactor and optimize Sage's Combinatorics (D5.11)

Progress report on other tasks

T5.1: Pari

T5.2: GAP

T5.4: Singular

## Task 5.6: HPC infrastructure for Combinatorics

*Parallel algebraic computations exhibit high degrees of irregularity, at multiple levels and make no use of floating-point operations. This combination means that symbolic computations are not suited to conventional HPC paradigms.*

– P. Maier, D. Livesey, H.-W. Loidl, P. Trinder

### General situation

Given a set  $S$  of combinatorial objects we want to

- ▶ find (or pick at random) an element of  $S$
- ▶ list (or count) the elements of  $S$

### Specificities of combinatorics

- ▶ **huge**:  $S$  does not fit in memory
- ▶ **embarrassingly parallel**:  $S = S_1 \cup \dots \cup S_k$
- ▶ **unbalanced**:  $S_i$  sizes are highly unbalanced

## D5.11: Refactor and optimise the existing combinatorics Sage code using the new developed Pythran and Cython features.


Algorithmic innovations and software integration through examples:

1. Example 1: polytopes
2. Example 2: numerical semigroups

# Example 1: polytopes (description)

## Example

Some combinatorial objects can be encoded as integer vectors in polytopes.



- ▶ Efficient algorithms available (e.g. Double Description, Barvinok)
- ▶ High performance libraries (e.g. PPL, Normaliz, Polymake)
- ▶ Useful for many combinatorial problems

# Example 1: polytopes (ODK work)

## ODK outcomes

**pplpy library:** Python interface to the high performance Parma Polyhedra Library (rational polytopes)

**e-antic library:** C/C++ library for computations over embedded number fields. Building block for multithread polytope computations in Normaliz over number fields.

**Workshop:** Sage Days 84

**SageMath:** improvement of native code and libraries interfaces



# Example 1: polytopes (ODK work)

## ODK outcomes

**pplpy library:** Python interface to the high performance Parma Polyhedra Library (rational polytopes) **6 releases**

**e-antic library:** C/C++ library for computations over embedded number fields. Building block for multithread polytope computations in Normaliz over number fields. **release planned**


**Workshop:** Sage Days 84

**SageMath:** improvement of native code and libraries interfaces **10 tickets**

## Example 2: semigroups (description)

### Example

Numerical semigroup are certain subsets of non-negative integers.



- ▶ Many mathematical open mathematical questions
- ▶ Highly unbalanced tree

## Example 2: semigroups (ODK work)

### A SageMath implementation of work-stealing map-reduce

- ▶ Work stealing algorithm (Leiserson-Blumofe)
- ▶ Easy to use, easy to call from SageMath with many use cases
- ▶ Scale well with the number of CPU cores and reasonably efficient (given that it is Python code).

A typical speedup (parallelization of binary sequences)

# processors	1	2	4	8
Speedup	$\times 1$	$\times 1.55$	$\times 2.43$	$\times 2.87$

## Example 2: semigroups (ODK work)

### A SageMath implementation of work-stealing map-reduce

- ▶ Work stealing algorithm (Leiserson-Blumofe)
- ▶ Easy to use, easy to call from SageMath with many use cases
- ▶ Scale well with the number of CPU cores and reasonably efficient (given that it is Python code).

A typical speedup (parallelization of binary sequences)

# processors	1	2	4	8
Speedup	×1	×1.55	×2.43	×2.87

integrated in SageMath

## Example 2: semigroups (ODK work)

### HPCombi: an experimental C++ library

- ▶ innovative SIMD code for combinatorics
- ▶ fine grained parallelization using Intel Cilk technology
- ▶ already in use for mathematical projects (numerical semigroups)

Operation	Speedup
Sum of a vector of bytes	×3.81
Sorting a vector of bytes	×21.3
Inverting a permutation	×1.97
Number of cycles of a permutation	×41.5
Number of inversions of a permutation	×9.39
Cycle type of a permutation	×8.94

## Example 2: semigroups (ODK work)

### HPCombi: an experimental C++ library

- ▶ innovative SIMD code for combinatorics
- ▶ fine grained parallelization using Intel Cilk technology
- ▶ already in use for mathematical projects (numerical semigroups)

Operation	Speedup
Sum of a vector of bytes	×3.81
Sorting a vector of bytes	×21.3
Inverting a permutation	×1.97
Number of cycles of a permutation	×41.5
Number of inversions of a permutation	×9.39
Cycle type of a permutation	×8.94

Open-source library



# Outline

## Workpackage management

## Exact linear algebra

Exact linear algebra algorithms and implementations (D5.12).

Distributed computing

## Combinatorics

Refactor and optimize Sage's Combinatorics (D5.11)

## Progress report on other tasks

T5.1: Pari

T5.2: GAP

T5.4: Singular

## D5.16: Pari Suite release, fully supporting parallelization

- ▶ D5.10 (merged in D5.16):  
Generic parallelization engine is now mature (released since nov.2016).  
Support POSIX-threads and MPI.
- ▶ Current work: applying it throughout the library
  - ▶ Chinese remaindering
  - ▶ Rational linear algebra
  - ▶ Discrete logarithm
  - ▶ Resultants
  - ▶ APRCL primality testing



### D5.15: Final report of GAP development

- ▶ 8 releases were cut integrating contributions of D3.11 and D5.15
- ▶ Towards an integration of HPC-GAP: main release GAP-4.9
  - ▶ Build system refactoring
  - ▶ Ability to compile in HPC-GAP compatibility mode
- ▶ Work in progress:
  - ▶ Multithreaded linear algebra: at the level of the Meataxe library
  - ▶ Introspection functionalities: on-the-fly optimisation decision

### D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

## T5.4 Singular

### D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

#### **Parallelization of the (sparse) multiplication**

threads	time (ms)	speedup
1	148661	x1.0
2	76881	x1.9
3	54798	x2.7
4	42855	x3.4
5	37017	x4.0
6	30892	x4.8
7	28365	x5.2
8	28048	x5.3

### D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

#### **Parallelization of the (sparse) multiplication**

threads	time (ms)	speedup
1	148661	x1.0
2	76881	x1.9
3	54798	x2.7
4	42855	x3.4
5	37017	x4.0
6	30892	x4.8
7	28365	x5.2
8	28048	x5.3

- ▶ Planned improvements to the memory manager  $\Rightarrow$  closer to linear scaling
- ▶ Parallel division and GCD implementations are in progress.
- ▶ Integration into Factory/Singular remains to be done

## Outcome of WorkPackage 5

	Review 1	Review 2	Final review
T5.1 Pari/GP			D5.16
T5.2 GAP			D5.15
T5.3 LinBox		D5.12	D5.14
T5.4 Singular	D5.6, D5.7		D5.13
T5.5 MPIR	D5.5, D5.7		
T5.6 Combinatorics	D5.1	D5.11	
T5.7 Pythran	D5.2	D5.11	
T5.8 SunGrid Engine	D5.3		

### Overall

- ▶ over 20 software releases were cut
- ▶ 7 research papers

Backup Slides

# Outsourced computing security (D5.12)

## Exploratory aspect: Computation over the Cloud

### Outsourcing computation on the cloud:

- ▶ trusted lightweight client computer
  - ▶ untrusted powerful cloud server
- ⇒ need for certification protocols

### Multiparty computation:

- ▶ each player contribute with a share of the input
- ▶ shares must remain private

# Outsourced computing security (D5.12)

## Exploratory aspect: Computation over the Cloud

### Outsourcing computation on the cloud:

- ▶ trusted lightweight client computer
  - ▶ untrusted powerful cloud server
- ⇒ need for certification protocols

### Multiparty computation:

- ▶ each player contribute with a share of the input
- ▶ shares must remain private

## Contribution


ISSAC'17: Linear time certificates for LU, Det., Rank Profile matrix, etc

In submission: Secure multiparty Strassen's algorithm



# Exploration: Python, Cython, Pythran and Cilk (D5.8)

*workshop: Interfacing (math) software with low level libraries*



- ▶ Python (interpreted language) slow.
- ▶ Cython/Pythran (Python to C compilers) efficient.
- ▶ Cilk (SIMD and multithread parallelism) very promising.