



Workpackage 5: High Performance Mathematical Computing

Second OpenDreamKit Project review

Luxembourg, October 30, 2018

High performance mathematical computing

Mathematical computing

Computing with a large variety of objects

► $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

for applications where all digits matter.

Mathematical computing

Computing with a large variety of objects

► $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

► Polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

for applications where all digits matter.

Mathematical computing

Computing with a large variety of objects

- ▶ $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$
$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

for applications where all digits matter.

Mathematical computing

Computing with a large variety of objects

- ▶ $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

for applications where all digits matter.

Mathematical computing

Computing with a large variety of objects

- ▶ $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$
- ▶ Matrices of polynomials over $\mathbb{Z}, \mathbb{Q}, \mathbb{Z}/p\mathbb{Z}, \mathbb{F}_q,$

17541718814389012164632

$$\frac{2}{5}x^3 + x^2 - \frac{1}{19}x + 2$$

$$\begin{bmatrix} 27 & 3 & -1 \\ 9 & 0 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 3x^2+3 & 2x^2+3 \\ 4x^2+1 & x^2+4x+4 \end{bmatrix}$$

- ▶ Tree algebra

for applications where all digits matter.

$$\frac{3q^2 - q^5}{q^5 + 2q^4 + 3q^3 + 3q^2 + 2q + 1} \begin{array}{c} (a) \\ / \quad \backslash \\ (b) \quad (c) \\ \quad \quad | \\ \quad \quad (d) \end{array} + \frac{2q}{q^4 + q^3 + 2q^2 + q + 1} \begin{array}{c} (a) \\ | \\ (b) \\ / \quad \backslash \\ (c) \quad (d) \end{array}$$

High performance mathematical computing

Need for High performance: applications where size matters:

Experimental maths: testing conjectures

- larger instances give higher confidence

High performance mathematical computing

Need for High performance: applications where size matters:

Experimental maths: testing conjectures

- ▶ larger instances give higher confidence

Algebraic cryptanalysis: security = computational difficulty

- ▶ key size determined by the largest solvable problem

Example

Breaking RSA by integer factorization: $n = pq$. Last record:

- ▶ n of 768 bits
- ▶ linear algebra in dimension 192 796 550 over \mathbb{F}_2 (105Gb)
- ▶ About 2000 CPU years

High performance mathematical computing

Need for High performance: applications where size matters:

Experimental maths: testing conjectures

- ▶ larger instances give higher confidence

Algebraic cryptanalysis: security = computational difficulty

- ▶ key size determined by the largest solvable problem

Example

Breaking RSA by integer factorization: $n = pq$. Last record:

- ▶ n of 768 bits
- ▶ linear algebra in dimension 192 796 550 over \mathbb{F}_2 (105Gb)
- ▶ About 2000 CPU years

3D data analysis, shape recognition:

- ▶ via persistent homology
- ▶ large sparse matrices over \mathbb{F}_2, \mathbb{Z}

Goal: delivering high performance to maths users

Systems :

GAP

PARI/GP

SageMath

Singular

Components :

FLINT

MPIR

LinBox

PPL

NumPy

Goal: delivering high performance to maths users

Harnessing modern hardware \rightsquigarrow parallelisation

- ▶ in-core parallelism (SIMD vectorisation)
- ▶ multi-core parallelism
- ▶ distributed computing: clusters, cloud

Systems :

GAP

PARI/GP

SageMath

Singular

Components :

FLINT

MPIR

LinBox

PPL

NumPy

Architectures :

SIMD

**Multicore
server**

HPC cluster

Cloud

Goal: delivering high performance to maths users

Languages

- ▶ Computational Maths software uses high level languages (e.g. Python)
 - ▶ High performance delivered by languages close to the metal (C, assembly)
- ~> compilation, automated optimisation

Systems :

GAP

PARI/GP

SageMath

Singular

Components :

FLINT

MPIR

LinBox

PPL

NumPy

Languages :

Cython

Python

Pythran

C

Architectures :

SIMD

**Multicore
server**

HPC cluster

Cloud

High performance mathematical computing

Goal:

- ▶ Improve/Develop parallelization of software components
- ▶ Expose them through the software stack
- ▶ Offer High Performance Computing to VRE's users

Milestone M8: Seamless use of parallel computing architecture in the VRE (proof of concept)

Astrid wants to run compute intensive routines involving both dense linear algebra and combinatorics. She has access through a JupyterHub-based VRE to a high end multi-core machine which includes a vanilla SAGE installation. She automatically benefits from the HPC features of the underlying specialized libraries (LinBox, ...). This is a proof of concept of the overall framework to integrate the HPC advances of specialized libraries into a general purpose VRE. It will prepare the final integration of a broader set of such parallel features for the end of the project

Addressing recommendations of review 1




Recommendation 10: *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be nearer the leading edge.*

Addressing recommendations of review 1

Recommendation 10: *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be **nearer the leading edge**.*

Leading edge achievements in linear algebra

- ▶ symmetric factorization outperforms LAPACK implementation
- ▶ new non-hierarchical generator for quasiseparable matrices
- ▶ large scale parallelization of rational linear solver






Addressing recommendations of review 1

Recommendation 10: *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be nearer the leading edge.*

Leading edge achievements in linear algebra

- ▶ symmetric factorization outperforms LAPACK implementation
- ▶ new non-hierarchical generator for quasiseparable matrices
- ▶ large scale parallelization of rational linear solver



Interactions and contacts made:

- ▶ interaction with the BLIS group (vectorization and implementation of Strassen's algorithm)
- ▶ on-going collaboration with T. Mary (Mumps) and S. Chandrasekaran (UCSB) on quasiseparable matrix algorithmic

Outline

Deliverables under review for the period

D5.12: Exact linear algebra algorithms and implementations.

D5.11: Refactor and optimize Sage's Combinatorics

Progress report on other deliverables

T5.1: Pari

T5.2: GAP

T5.3: LinBox

T5.4: Singular

Workpackage management

Milestone M8

Strengthening interactions with numerical HPC

Outline

Deliverables under review for the period

D5.12: Exact linear algebra algorithms and implementations.

D5.11: Refactor and optimize Sage's Combinatorics

Progress report on other deliverables

T5.1: Pari

T5.2: GAP

T5.3: LinBox

T5.4: Singular

Workpackage management

Milestone M8

Strengthening interactions with numerical HPC

Task 5.3: LinBox, High performance exact linear algebra

Mathematics is the art of reducing any problem to linear algebra

– W. Stein

Linear algebra: a HPC building block

Similarly as in numerical HPC:

- ▶ central elementary problem to which others reduce to
- ▶ (rather) simple algorithmic
- ▶ high compute/memory intensity

Specificities

- ▶ Multiprecision arithmetic \Rightarrow lifting from finite precision (\mathbb{F}_p)
- ▶ Rank deficiency \Rightarrow unbalanced dynamic blocking
- ▶ Early adopter of subcubic matrix arithmetic \Rightarrow recursion

D5.12: Exact linear algebra algorithms and implementations. Library maintenance and close integration in mathematical software for LinBox library



1. Algorithmic innovations:

- 1.1 Rank deficient dense Gaussian elimination
- 1.2 Quasiseparable matrices
- 1.3 Outsourced computing security

2. Software releases and integration:

- 2.1 LinBox ecosystem: LinBox, fflas-ffpack, givaro
- 2.2 SageMath integration

Rank deficient dense Gaussian elimination

[JSC'17] Fast computation of the rank profile matrix and the generalized Bruhat decomposition.

- ▶ Connecting Rank Profile Matrix and row and column echelon forms
- ▶ $O(r^\omega + mn)$ probabilistic time
- ▶ generalization over arbitrary rings


Rank deficient dense Gaussian elimination

[JSC'17] Fast computation of the rank profile matrix and the generalized Bruhat decomposition.

- ▶ Connecting Rank Profile Matrix and row and column echelon forms
- ▶ $\tilde{O}(r^\omega + mn)$ probabilistic time
- ▶ generalization over arbitrary rings


[ISSAC'18] Symmetric triangular factorization

- ▶ First unconditional recursive algorithm
- ▶ Pivoting revealing the Rank Profile Matrix
- ▶ $O(n^2 r^{\omega-2})$ ($= 1/3 n^3$ with $\omega = 3, r = n$)



LAPACK vs FFPACK modulo 8 388 593


n	LAPACK		FFPACK	
	dgetrf (LU)	dsytrf (LDLT)	fgetrf (LU)	fsytrf (LDLT)
5000	2.01s	1.60s	3.90s	1.59s
10000	14.95s	11.98s	24.12s	10.90s



Matrices with low off-diagonal rank

[ISSAC'16, JSC'18] New compact representation and algorithms

- ▶ Connection with rank profile matrix
- ▶ Matches the best space complexities: $O(ns)$
- ▶ Reduction to matrix multiplication: $O(ns^{\omega-1})$ for products
- ▶ Flat representation (non hierarchical)




Matrices with low off-diagonal rank

[ISSAC'16, JSC'18] New compact representation and algorithms

- ▶ Connection with rank profile matrix
- ▶ Matches the best space complexities: $O(ns)$
- ▶ Reduction to matrix multiplication: $O(ns^{\omega-1})$ for products
- ▶ Flat representation (non hierarchical)

Follow-up: on-going collaboration with numerical HPC experts:

- ▶ S. Chandrasekaran (UCSB)
- ▶ T. Mary (U. Manchester, Mumps)



Outsourced computing security

Exploratory aspect: Computation over the Cloud

Outsourcing computation on the cloud:

- ▶ trusted lightweight client computer
 - ▶ untrusted powerful cloud server
- ⇒ need for certification protocols

Multiparty computation:

- ▶ each player contribute with a share of the input
- ▶ shares must remain private

Outsourced computing security

Exploratory aspect: Computation over the Cloud

Outsourcing computation on the cloud:

- ▶ trusted lightweight client computer
 - ▶ untrusted powerful cloud server
- ⇒ need for certification protocols

Multiparty computation:

- ▶ each player contribute with a share of the input
- ▶ shares must remain private

Contribution

ISSAC'17: Linear time certificates for LU, Det., Rank Profile matrix, etc

In submission: Secure multiparty Strassen's algorithm

Software releases and integration

LinBox ecosystem

`givaro`: field/ring arithmetic

`fflas-ffpack`: dense linear algebra over finite field

`LinBox`: exact linear algebra

Tightly integrated in SageMath

Software releases and integration

LinBox ecosystem

<code>givaro</code> : field/ring arithmetic	4 releases
<code>fflas-ffpack</code> : dense linear algebra over finite field	6 releases
<code>LinBox</code> : exact linear algebra	6 releases
Tightly integrated in SageMath	13 tickets

Software releases and integration

LinBox ecosystem

<code>givaro</code> : field/ring arithmetic	4 releases
<code>fflas-ffpack</code> : dense linear algebra over finite field	6 releases
<code>LinBox</code> : exact linear algebra	6 releases
Tightly integrated in SageMath	13 tickets

Featuring

- ▶ Full functional implementations of new algorithmic contributions
- ▶ Improved vectorization and parallel routines
- ▶ Drastic improvement of reliability (continuous integration, test-suite coverage, randomized certificates, etc)

Task 5.6: HPC infrastructure for Combinatorics

General situation

Given a set S of combinatorial objects we want to

- ▶ test conjectures: e.g. find an element of S satisfying a certain property
- ▶ count or list the elements of S having this property

Specificities of combinatorics

Typically

- ▶ **huge**: S does not fit in memory
- ▶ **embarrassingly parallel**: $S = S_1 \cup \dots \cup S_k$
- ▶ **unbalanced**: S_i sizes are highly unbalanced

D5.11: Refactor and optimise the existing combinatorics Sage code using the new developed Pythran and Cython features.




1. Algorithmic innovations, software integration and experimentations through examples
2. Software releases and integration:

Concrete mathematical problem 1: polytopes

Example

Some combinatorial objects can be encoded as integer vectors in polytopes.



- ▶ Efficient algorithms available (Double description, Barvinok)
- ▶ High performance libraries implementing them PPL, Normaliz
- ▶ Useful for many combinatorial applications


ODK outcomes

- ▶ `pp1py` library: Python interface to the high performance Parma Polyhedra Library (rational polytopes)
- ▶ `e-antic` library: C/C++ library for computations over embedded number fields. Building block for multithread polytope computations in Normaliz over number fields.
- ▶ *Sage Days 84* workshop

Concrete mathematical problem 2: semigroups

Example

Numerical semigroup are certain subsets of non-negative integers.



- ▶ Many mathematical open mathematical questions.
- ▶ Highly unbalanced tree.

A Python implementation

- ▶ Work stealing algorithm (Leiserson-Blumofe)
- ▶ Easy to use, easy to call from SageMath with many use cases
- ▶ Scale well with the number of CPU cores and reasonably efficient (given that it is Python code).

A typical speedup obtained for binary sequences

# processors	1	2	4	8
Time (s)	250	161	103	87


Low-level parallelization for small combinatorial objects

HPCombi: an experimental C++ library

- ▶ innovative SIMD code for combinatorics
- ▶ fine grained parallelization using Intel Cilk technology
- ▶ already in use for mathematical projects (numerical semigroups)

Operation	Speedup
Sum of a vector of bytes	×3.81
Sorting a vector of bytes	×21.3
Inverting a permutation	×1.97
Number of cycles of a permutation	×41.5
Number of inversions of a permutation	×9.39
Cycle type of a permutation	×8.94

Exploration: Python, Cython, Pythran and Cilk (D5.8)



- ▶ Python (interpreted language) slow.
- ▶ Cython/Pythran (Python to C compilers) efficient.
- ▶ Cilk (SIMD and multithread parallelism) very promising.

Software releases and integration

Libraries releases and integration

SageMath: improvement of native code and libraries interfaces

HPCombi: C++ library for small combinatorial objects using SIMD instructions and fine grained parallelization (Cilk)

e-antic : C/C++ library for (arbitrary precision) embedded number field computations

pplpy: Python library interface to PPL library on Polytope

Software releases and integration

Libraries releases and integration

- SageMath:** improvement of native code and libraries interfaces **5 tickets**
- HPCombi:** C++ library for small combinatorial objects using SIMD instructions and fine grained parallelization (Cilk) **experimental**
- e-antic** : C/C++ library for (arbitrary precision) embedded number field computations **planned release**
- pplpy:** Python library interface to PPL library on Polytope **6 releases**

Software releases and integration

Libraries releases and integration

- SageMath:** improvement of native code and libraries interfaces **5 tickets**
- HPCombi:** C++ library for small combinatorial objects using SIMD instructions and fine grained parallelization (Cilk) **experimental**
- e-antic :** C/C++ library for (arbitrary precision) embedded number field computations **planned release**
- pplpy:** Python library interface to PPL library on Polytope **6 releases**

Featuring

- ▶ SIMD instructions for combinatorial enumeration
- ▶ Generic map-reduce in SageMath
- ▶ Faster and more parallel code inside SageMath

Outline

Deliverables under review for the period

D5.12: Exact linear algebra algorithms and implementations

D5.11: Refactor and optimize Sage's Combinatorics

Progress report on other deliverables

T5.1: Pari

T5.2: GAP

T5.3: LinBox

T5.4: Singular

Workpackage management

Milestone M8

Strengthening interactions with numerical HPC

D5.16: Pari Suite release, fully supporting parallelization

- ▶ D5.10 (merged in D5.16):
Generic parallelization engine is now mature (released since nov.2016).
Support POSIX-threads and MPI.
- ▶ Current work: applying it throughout the library
 - ▶ Chinese remaindering
 - ▶ Rational linear algebra
 - ▶ Discrete logarithm
 - ▶ Resultants
 - ▶ APRCL primality testing

D5.15: Final report of GAP development



- ▶ 8 releases were cut integrating contributions of D3.11 and D5.15
- ▶ Towards an integration of HPC-GAP: main release GAP-4.9
 - ▶ Build system refactoring
 - ▶ Ability to compile in HPC-GAP compatibility mode
- ▶ Work in progress:
 - ▶ Multithreaded linear algebra: at the level of the Meataxe library
 - ▶ Introspection functionalities: on-the-fly optimisation decision

D5.14: Distributed exact linear system solving

- ▶ 2 full time engineers
- ▶ Communication and serialization layer done
- ▶ Prototype MPI parallelization of Chinese remainder based solver.

D5.14: Distributed exact linear system solving

- ▶ 2 full time engineers
- ▶ Communication and serialization layer done
- ▶ Prototype MPI parallelization of Chinese remainder based solver.





T5.3 LinBox

D5.14: Distributed exact linear system solving

Roadmap:

- ▶ Major refactorization of LinBox solver code under way
- ▶ Parallelization of Dixon-lifting solver
- ▶ Hybrid combination of CRT+Dixon.
- ▶ Hybrid OpenMP-MPI implementation



T5.4 Singular

D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

T5.4 Singular

D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

Parallelization of the (sparse) multiplication

threads	time (ms)	speedup
1	148661	x1.0
2	76881	x1.9
3	54798	x2.7
4	42855	x3.4
5	37017	x4.0
6	30892	x4.8
7	28365	x5.2
8	28048	x5.3

D5.13: Parallel sparse polynomial multiplication

FLINT now supports fast sparse multivariate polynomials:

- ▶ addition, subtraction, multiplication,
- ▶ division, division with remainder, GCD
- ▶ evaluation (and partial evaluation), composition

Parallelization of the (sparse) multiplication

threads	time (ms)	speedup
1	148661	x1.0
2	76881	x1.9
3	54798	x2.7
4	42855	x3.4
5	37017	x4.0
6	30892	x4.8
7	28365	x5.2
8	28048	x5.3

- ▶ Planned improvements to the memory manager \Rightarrow closer to linear scaling
- ▶ Parallel division and GCD implementations are in progress.
- ▶ Integration into Factory/Singular remains to be done

Outline

Deliverables under review for the period

D5.12: Exact linear algebra algorithms and implementations

D5.11: Refactor and optimize Sage's Combinatorics

Progress report on other deliverables

T5.1: Pari

T5.2: GAP

T5.3: LinBox

T5.4: Singular

Workpackage management

Milestone M8




Strengthening interactions with numerical HPC

Recommendation 10: *Regarding WP5, make **contacts** with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be **nearer the leading edge**.*

Recommendation 10: *Regarding WP5, make contacts with HPC community in order to ascertain current state-of-the-art. The work in this WP needs to be **nearer the leading edge**.*

Leading edge achievements in linear algebra

- ▶ symmetric factorization outperforms LAPACK implementation
- ▶ new non-hierarchical generator for quasiseparable matrices
- ▶ large scale parallelization of rational linear solver



Strengthening interactions with numerical HPC community

Existing connection

- ▶ Dense linear algebra: numerical BLAS used for exact FFLAS for 17 years
- ▶ Pointwise interactions with J. Dongarra, L. Grigori, J-Y. L'Excellent, etc
- ▶ Publishing in major HPC venues: SIAM-PPSC, EuroPar, PMAA, ParCo
- ▶ Involvement in the French CNRS GDR-Calcul working group (Sci Comp)

Strengthening interactions with numerical HPC community

Existing connection

- ▶ Dense linear algebra: numerical BLAS used for exact FFLAS for 17 years
- ▶ Pointwise interactions with J. Dongarra, L. Grigori, J-Y. L'Excellent, etc
- ▶ Publishing in major HPC venues: SIAM-PPSC, EuroPar, PMAA, ParCo
- ▶ Involvement in the French CNRS GDR-Calcul working group (Sci Comp)

Recently established

- ▶ With the BLIS group:
 - ▶ Report SIMD vectorization bug and share user experience
 - ▶ Implementation of Strassen's algorithm
- ▶ On-going collaboration with T. Mary (Mumps) and S. Chandrasekaran (UCSB) on quasiseparable matrix algorithmic