

# Symmetric Indefinite Triangular Factorization Revealing the Rank Profile Matrix\*

Jean-Guillaume Dumas

Université Grenoble Alpes

Laboratoire Jean Kuntzmann, CNRS, UMR 5224

38058 Grenoble, CEDEX 9, France

Jean-Guillaume.Dumas@univ-grenoble-alpes.fr

Clément Pernet

Université Grenoble Alpes

Laboratoire Jean Kuntzmann, CNRS, UMR 5224

38058 Grenoble, CEDEX 9, France

Clement.Pernet@univ-grenoble-alpes.fr

## ABSTRACT

We present a novel recursive algorithm for reducing a symmetric matrix to a triangular factorization which reveals the rank profile matrix. That is, the algorithm computes a factorization  $\mathbf{P}^T \mathbf{A} \mathbf{P} = \mathbf{LDL}^T$  where  $\mathbf{P}$  is a permutation matrix,  $\mathbf{L}$  is lower triangular with a unit diagonal and  $\mathbf{D}$  is symmetric block diagonal with  $1 \times 1$  and  $2 \times 2$  antidiagonal blocks. This algorithm requires  $O(n^2 r^{\omega-2})$  arithmetic operations, with  $n$  the dimension of the matrix,  $r$  its rank and  $\omega$  an admissible exponent for matrix multiplication. Furthermore, experimental results demonstrate that our algorithm has very good performance: its computational speed matches that of its numerical counterpart and is twice as fast as the unsymmetric exact Gaussian factorization. By adapting the pivoting strategy developed in the unsymmetric case, we show how to recover the rank profile matrix from the permutation matrix and the support of the block-diagonal matrix. We also note that there is an obstruction in characteristic 2 for revealing the rank profile matrix, which requires to relax the shape of the block diagonal by allowing the 2-dimensional blocks to have a non-zero bottom-right coefficient. This relaxed decomposition can then be transformed into a standard  $\mathbf{PLDL}^T \mathbf{P}^T$  decomposition at a negligible cost.

## ACM Reference Format:

Jean-Guillaume Dumas and Clément Pernet. 2018. Symmetric Indefinite Triangular Factorization Revealing the Rank Profile Matrix. In *ISSAC '18: 2018 ACM International Symposium on Symbolic and Algebraic Computation, July 16–19, 2018, New York, NY, USA*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3208976.3209019>

## 1 INTRODUCTION

Computing a triangular factorization of a symmetric matrix is a commonly used routine to solve symmetric linear systems, or to compute the signature of symmetric bilinear forms. Besides the fact that it is expected to save half of the arithmetic cost of a standard (non-symmetric) Gaussian elimination, it can also recover invariants, such as the signature, specific to symmetric matrices, and thus, e.g., be used to certify positive or negative definiteness or semidefiniteness [13, Corollary 1].

\*This work is partly funded by the OpenDreamKit Horizon 2020 European Research Infrastructures project (#676541).

Publication rights licensed to ACM. ACM acknowledges that this contribution was authored or co-authored by an employee, contractor or affiliate of a national government. As such, the Government retains a nonexclusive, royalty-free right to publish or reproduce this article, or to allow others to do so, for Government purposes only.

ISSAC '18, July 16–19, 2018, New York, NY, USA

© 2018 Copyright held by the owner/author(s). Publication rights licensed to the Association for Computing Machinery.

ACM ISBN 978-1-4503-5550-6/18/07...\$15.00

<https://doi.org/10.1145/3208976.3209019>

It is a fundamental computation in numerical linear algebra, and is therefore most often presented in the setting of real matrices. When the matrix is positive definite, the Cholesky factorization can be defined:  $\mathbf{A} = \mathbf{LL}^T$ , where  $\mathbf{L}$  is lower triangular for which square roots of diagonal elements have to be extracted. Alternatively, gathering the diagonal elements in a central diagonal matrix yields the LDLT factorization  $\mathbf{A} = \mathbf{LDL}^T$  which no longer requires square roots. Similarly as for the LU decomposition, it is only defined for matrices with generic rank profile, i.e. having their  $r = \text{rank}(\mathbf{A})$  first leading principal minors non-zero. For arbitrary matrices, symmetric permutations may lead to the former situations:  $\mathbf{PAP}^T = \mathbf{LDL}^T$ . However, this is unfortunately not always the case. For instance there is no permutation  $\mathbf{P}$  such that  $\mathbf{P}^T \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \mathbf{P}$  has an LDLT factorization with a diagonal  $\mathbf{D}$ . This led to a series of generalizations where the matrix  $\mathbf{D}$  was replaced first by a tridiagonal symmetric matrix by Parlett and Reid [14], improved by Aasen [1], achieving half the arithmetic cost of Gaussian elimination. Bunch and Kaufman then replaced this tridiagonal matrix by a block diagonal composed of 1 or 2-dimensional antidiagonal blocks to obtain an *indefinite triangular factorization*.

*Pivoting.* In numerical linear algebra, the choice of the permutation matrix is mainly driven by the need to ensure a good numerical quality of the decomposition. Bunch and Parlett [6] use a full pivoting technique, requiring a cubic number of tests. Bunch and Kaufman's pivoting strategy, implemented in LAPACK, uses a partial pivoting requiring only a quadratic number of tests.

In the context of exact linear algebra, for instance when computing over a finite field, numerical stability is no longer an issue. However, the computation of echelon forms and rank profiles, central in many applications, impose further constraints on the pivoting. A characterization of the requirements for the pivoting strategy is given in [10, 11] so that a PLUQ decomposition can reveal these rank profiles and echelon forms in the non-symmetric case. In particular, it is shown that pivot selection minimizing the lexicographic order on the coordinate of the pivot, combined with row and column rotations to move the pivot to the diagonal, enable the computation of the rank profile matrix, an invariant from which all rank profile information or the row and the column echelon form can be recovered.

*Recursive algorithms.* As in numerical linear algebra, we try to gather arithmetic operations in level 3 BLAS operations (matrix multiplication based), for it delivers the best computation throughput. Numerical software often use tiled implementations, especially when the pivoting is more constrained by the symmetry [12, 16], or in order to define communication avoiding variants [4]. In exact

linear algebra sub-cubic matrix multiplication, such as Strassen's algorithm, can be extensively used with no numerical instability issues. This led to the design of recursive algorithms, which was proven successful in the unsymmetric case, including for shared memory parallel computations [9].

**Contribution.** The contribution here is to propose a recursive algorithm producing a symmetric factorization PLDLTPT over any field, from which the rank profile matrix of the input can be recovered. This algorithm is a recursive variant of Bunch and Kaufman's algorithm [5] where the pivoting strategy has been replaced by the one developed previously by the authors in the unsymmetric case [11]. Compared to the recursive adaptation of Aasen's algorithm in [15], our algorithm leads to a similar data partitioning but does not suffer from an arithmetic overhead with respect to Aasen's algorithm. Our algorithm has time complexity  $O(n^2 r^{\omega-2})$  where  $\omega$  is an admissible exponent for matrix multiplication and  $r$  is the rank of the input matrix of dimension  $n$ . With  $\omega = 3$ , the leading constant in the time complexity is  $1/3$ , matching that of the best alternative algorithms based on cubic time linear algebra.

In Section 2 we show that in characteristic two the rank profile matrix can not always be revealed by a symmetric factorization with antidiagonal blocks: sometimes antitriangular blocks are also required. Then we recall in Section 3 the main required level 3 linear algebra subroutines. In Section 4 we present the main recursive algorithm. An alternative iterative Crout variant is presented in Section 5 to be used as a base case in the recursion. We finally show, in Section 6, experiments of the resulting implementation over a finite field. They demonstrate the efficiency of cascading the recursive algorithm with the base case variant, especially with matrices involving a lot of pivoting. They finally confirm a speed-up by a factor of about 2 compared to the state of the art unsymmetric Gaussian elimination.

## 2 THE SYMMETRIC RANK PROFILE MATRIX

### 2.1 The pivoting matrix

Theorem 1 recalls the definition of the rank profile matrix (RPM).

**THEOREM 1 ([10]).** *Let  $A \in \mathbb{F}^{m \times n}$ . There exists a unique  $m \times n$   $\{0, 1\}$ -matrix  $\mathcal{R}_A$  with  $r$  1's in rook placement of which every leading sub-matrix has the same rank as the corresponding leading sub-matrix of  $A$ . This matrix is called the rank profile matrix of  $A$ .*

**LEMMA 1.** *A symmetric matrix has a symmetric rank profile matrix.*

**PROOF.** Otherwise, the rank of some leading submatrix of  $A$  and the same leading submatrix of  $A^T$  would be different which is absurd.  $\square$

Also, any symmetric matrix has a triangular decomposition  $A = \text{PLDL}^T \mathbf{P}^T$  where  $L$  is unit lower triangular,  $P$  is a permutation matrix and  $D$  is block diagonal, formed by 1-dimensional scalar blocks or 2-dimensional antidiagonal blocks of the form  $\begin{bmatrix} 0 & x \\ x & 0 \end{bmatrix}$ . We further define the support matrix of such a block diagonal matrix as the localization of the non-zero elements:

**DEFINITION 1.** *The support matrix of a block diagonal matrix  $D$ , formed by 1-dimensional scalar blocks or 2-dimensional blocks*

*of the form  $\begin{bmatrix} 0 & x \\ x & 0 \end{bmatrix}$  is the block diagonal  $\{0, 1\}$ -matrix  $\Psi_D$  such that  $D = \Psi_D \bar{D}$ , with  $\bar{D}$  a diagonal matrix.*

**DEFINITION 2.** *The pivoting matrix of a  $\text{PLDL}^T \mathbf{P}^T$  decomposition is the matrix  $\Pi = \mathbf{P} \Psi_D \mathbf{P}^T$ .*

**DEFINITION 3.** *A  $\text{PLDL}^T \mathbf{P}^T$  decomposition is said to reveal the rank profile matrix of a symmetric matrix  $A$  if its pivoting matrix equals the rank profile matrix of  $A$ .*

### 2.2 Antitriangular blocks in characteristic two

In zero or odd characteristic, we show next that one can always find such a PLDLTPT decomposition revealing the rank profile matrix. In characteristic two, however, this is not always possible.

**LEMMA 2.** *In characteristic 2, there is no symmetric indefinite elimination revealing the rank profile matrix of  $A = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}$ .*

**PROOF.** Let  $J$  be the  $2 \times 2$  anti-diagonal identity matrix. This is also the rank profile matrix of  $A$ . Now, we let  $L = \begin{bmatrix} 1 & 0 \\ x & 1 \end{bmatrix}$ ,  $\bar{D} = \begin{bmatrix} y & 0 \\ 0 & z \end{bmatrix}$ . As the permutation matrices involved,  $P$  and  $\Psi_D$ , can only be either the identity matrix or  $J$ , there are then four cases:

- (1)  $A = L \cdot \bar{D} \cdot L^T = \begin{bmatrix} y & xy \\ xy & x^2y + z \end{bmatrix}$ , but  $y = 0$  and  $\text{rank}(\bar{D}) = \text{rank}(A) = 2$  are incompatible.
- (2)  $A = J \cdot L \cdot \bar{D} \cdot L^T \cdot J^T = \begin{bmatrix} x^2y + z & xy \\ xy & y \end{bmatrix}$ , but  $I = J \cdot I \cdot J \neq J = \mathcal{R}_A$ .
- (3)  $A = L \cdot \bar{D} \cdot J \cdot L^T = \begin{bmatrix} 0 & y \\ z & xy + xz \end{bmatrix}$ , but we need  $y = z$  for the symmetry and then  $2xy = 0 \neq 1$  in characteristic 2.
- (4)  $A = J \cdot L \cdot \bar{D} \cdot J \cdot L^T \cdot J^T = \begin{bmatrix} xy + xz & z \\ y & 0 \end{bmatrix}$  but the bottom right coefficient of  $A$  is non zero.

$\square$

However, one can generalize the PLDLTPT decomposition to a block diagonal matrix  $D$  having 2-dimensional blocks of the form  $\begin{bmatrix} 0 & c \\ c & d \end{bmatrix}$  (lower antitriangular). Then the notion of support matrix can be generalized: for such a block diagonal matrix  $D$ ,  $\Psi_D$  is the block diagonal  $\{0, 1\}$  matrix such that  $D = \Psi_D \bar{D}$ , with  $\bar{D}$  an upper triangular bidiagonal matrix (or equivalently such that  $D = \bar{D} \Psi_D$ , with  $\bar{D}$  lower triangular bidiagonal). For instance  $\begin{bmatrix} 0 & c \\ c & d \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} c & d \\ 0 & c \end{bmatrix} = \begin{bmatrix} c & 0 \\ d & c \end{bmatrix} \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$ .

With these generalized definitions, we show in Section 4, that there exists RPM-revealing PLDLTPT decompositions.

### 2.3 Antitriangular decomposition

Then, such a generalized decomposition can always be further reduced to a strict PLDLTPT decomposition by eliminating each of the antitriangular blocks. For this, the observation is that in characteristic two, a symmetric lower antitriangular  $2 \times 2$  block is invariant under any symmetric triangular transformation:  $\begin{bmatrix} 1 & \\ x & 1 \end{bmatrix} \begin{bmatrix} c \\ c & d \end{bmatrix} \begin{bmatrix} 1 & x \\ & 1 \end{bmatrix} =$

$$\begin{bmatrix} c \\ c & 2cx + d \end{bmatrix} \equiv \begin{bmatrix} c \\ c & d \end{bmatrix} \pmod{2} = \begin{bmatrix} 1 & \\ & 1 \end{bmatrix} \begin{bmatrix} c \\ c & d \end{bmatrix} \begin{bmatrix} 1 & \\ & 1 \end{bmatrix}. \text{ Thus for each}$$



or invertible. We will later give the general presentation of the algorithm where its rank could be arbitrary.

Let  $M \in \mathbb{F}^{(m+n) \times (m+n)}$  be the symmetric matrix to be factorized.

Consider its block decomposition  $M = \begin{bmatrix} A & B \\ B^T & C \end{bmatrix}$  where  $A \in \mathbb{F}^{m \times m}$  and  $C \in \mathbb{F}^{n \times n}$  are also symmetric.

If  $A$  is full rank, then a recursive call will produce  $A = PLDL^T P^T$ , and  $M$  can thus be decomposed as:

$$M = \begin{bmatrix} P & 0 \\ 0 & I \end{bmatrix} \begin{bmatrix} L & 0 \\ G & I \end{bmatrix} \begin{bmatrix} D & 0 \\ 0 & Z \end{bmatrix} \begin{bmatrix} L^T & G^T \\ 0 & I \end{bmatrix} \begin{bmatrix} P^T & 0 \\ 0 & I \end{bmatrix},$$

where  $G$  is such that  $PLDG^T = B$  and  $Z = C - GDG^T$ . Thus  $G$  can be computed as the transpose of  $D^{-1}L^{-1}P^{-1}B$  which can be obtained by a call to `trsm`, some permutations and a diagonal scaling. Then  $Z$  is computed by a call to `syrdk`. A second recursive call will then decompose  $Z$  and lead to the final factorization of  $M$ .

Now if  $A$  is the zero matrix, one is reduced to factorize the matrix  $M = \begin{bmatrix} 0 & B \\ B^T & C \end{bmatrix}$ . In order to recover the rank profile matrix, one has to first look for pivots in  $B$  before considering the block  $C$ . Therefore diagonal pivoting is not an option here. Then the matrix  $B$ , which we assume has full rank for the moment, can be decomposed in a  $PLDUQ$  factorization ( $P$  and  $Q$  are permutation matrices,  $L$  and  $U$  are respectively unit lower and unit upper triangular,  $D$  is diagonal). We then need to distinguish two cases depending on whether the field characteristic is two or not.

**4.1.1 Zero or odd characteristic case.** If the characteristic is zero or odd,  $M$  can thus be decomposed as:

$$M = \begin{bmatrix} P & 0 \\ 0 & Q^T \end{bmatrix} \begin{bmatrix} L & 0 \\ G & U^T \end{bmatrix} \begin{bmatrix} 0 & D \\ D & 0 \end{bmatrix} \begin{bmatrix} L^T & G^T \\ 0 & U \end{bmatrix} \begin{bmatrix} P^T & 0 \\ 0 & Q \end{bmatrix},$$

where  $G$  is such that  $Q^T(GDU + U^T DG^T)Q = C$ . To compute  $G$ , one can first permute  $C$  to get  $C' = QCQ^T$  (which remains symmetric) and then use a call to `trssyr2k`.

**4.1.2 Characteristic two case.** In characteristic two, the equation  $GDU + U^T DG^T = QCQ^T$  in unknown  $G$  has in general no solution (as soon as  $C$  has a non-zero diagonal element).

However, one can still relax Problem 1 and allow the elimination to leave a diagonal of elements not zeroed out. Following Lemma 2, the idea is then to decompose  $M$  into a block tridiagonal form:

$$M = \begin{bmatrix} P & 0 \\ 0 & Q^T \end{bmatrix} \begin{bmatrix} L & 0 \\ G & U^T \end{bmatrix} \begin{bmatrix} 0 & D \\ D & \Delta \end{bmatrix} \begin{bmatrix} L^T & G^T \\ 0 & U \end{bmatrix} \begin{bmatrix} P^T & 0 \\ 0 & Q \end{bmatrix},$$

where  $\Delta$  is a diagonal matrix and now  $G$  is such that  $Q^T(GDU + U^T DG^T + U^T \Delta U)Q = C$ . Therefore  $\Delta$  can be chosen such that the diagonal of  $C'' = QCQ^T - U^T \Delta U = C' - U^T \Delta U$  is zero. As  $U$  is unit upper triangular, a simple pass over its coefficients is sufficient to find such a  $\Delta$ : let  $\Delta_{ii} = C'_{ii} - \sum_{j=1}^{i-1} \Delta_{jj} U_{j,i}^2$ . The algorithm is thus to permute  $C$  to get  $C'$ ; then compute  $\Delta$  with the recursive relation above and update  $C'' = C' - U^T \Delta U$  with a `syrdk`.  $C''$  remains symmetric but with a zero diagonal and now `trssyr2k` can be applied.

## 4.2 The actual recursive algorithm

**4.2.1 First phase: recursive elimination.** In the general case, the leading matrices are not full rank, and we have to consider intermediate steps. For the symmetric matrix  $M \in \mathbb{F}^{(m+n) \times (m+n)}$  of Section 4.1, its leading principal block  $A \in \mathbb{F}^{m \times m}$  has rank  $r \leq m$ . Thus its actual recursive decomposition is of the form:

$$A = P_1 \begin{bmatrix} L_1 \\ M_1 \end{bmatrix} D_1 \begin{bmatrix} L_1^T & M_1^T \end{bmatrix} P_1^T,$$

where  $L_1 \in \mathbb{F}^{r \times r}$  is full rank unit lower triangular,  $D_1 \in \mathbb{F}^{r \times r}$  is block diagonal with 1 or 2-dimensional diagonal blocks, and  $M_1 \in \mathbb{F}^{(m-r) \times r}$ . Therefore, forgetting briefly the permutations, the decomposition of  $M$  becomes:

$$M = \begin{bmatrix} L_1 & 0 & 0 \\ M_1 & I & 0 \\ G & 0 & I \end{bmatrix} \begin{bmatrix} D_1 & 0 & 0 \\ 0 & 0 & Y \\ 0 & Y^T & Z \end{bmatrix} \begin{bmatrix} L_1^T & M_1^T & G^T \\ 0 & I & 0 \\ 0 & 0 & I \end{bmatrix},$$

where  $Y$  is such that  $B = \begin{bmatrix} L_1 \\ M_1 \end{bmatrix} D_1 G^T + \begin{bmatrix} 0 \\ Y \end{bmatrix}$ .

From this point on, there remains to factorize the submatrix  $\begin{bmatrix} 0 & Y \\ Y^T & Z \end{bmatrix}$ . This will be carried out by the algorithm described in the next section, working on a matrix with a zero leading principal submatrix. Supposing for now that this is possible, Algorithm 2 summarizes the whole procedure.

---

### Algorithm 2 Recursive symmetric indefinite elimination

---

**Require:**  $A \in \mathbb{F}^{m \times m}$  and  $C \in \mathbb{F}^{n \times n}$  both symmetric,  $B \in \mathbb{F}^{m \times n}$ .  
**Ensure:**  $P$  permutation,  $L$  unit lower triangular,  $D$  block diagonal,  
s.t.  $\begin{bmatrix} A & B \\ B^T & C \end{bmatrix} = PLDL^T P^T$ .

- 1: Decompose  $A = P_1 \begin{bmatrix} L_1 \\ M_1 \end{bmatrix} D_1 \begin{bmatrix} L_1^T & M_1^T \end{bmatrix} P_1^T$  {Alg. 2}
  - 2: let  $r = \text{rank}(A)$  s.t.  $L_1$  and  $D_1$  are  $r \times r$ .
  - 3:  $B' = P_1^T B$  {perm ( $P_1^T, B$ )}
  - 4: Split  $B' = \begin{bmatrix} B'_1 \\ B'_2 \end{bmatrix}$  where  $B'_1$  is  $r \times n$ .
  - 5:  $X \leftarrow L_1^{-1} B'_1$  {trsm ( $L_1, B'_1$ )}
  - 6:  $Y \leftarrow B'_2 - M_1 X$  {gemm ( $B'_2, M_1, X$ )}
  - 7:  $G \leftarrow X^T D_1^{-1}$  {scal ( $X^T, D_1^{-1}$ )}
  - 8:  $Z \leftarrow C - GD_1 G^T$  {syrdk ( $C, G, D_1$ )}
  - 9: Decompose  $\begin{bmatrix} 0 & Y \\ Y^T & Z \end{bmatrix} = P_2 L_2 D_2 L_2^T P_2^T$  {Alg. 3}
  - 10:  $P \leftarrow \begin{bmatrix} P_1 & 0 \\ 0 & I_n \end{bmatrix} \cdot \begin{bmatrix} I_r & 0 \\ 0 & P_2 \end{bmatrix}$
  - 11:  $N_1 \leftarrow P_2^T \begin{bmatrix} M_1 \\ G \end{bmatrix}$  {perm ( $P_2^T, \begin{bmatrix} M_1 \\ G \end{bmatrix}$ )}
  - 12:  $L \leftarrow \begin{bmatrix} L_1 & & \\ & N_1 & L_2 \\ & & \end{bmatrix}$
  - 13:  $D \leftarrow \begin{bmatrix} D_1 & & \\ & D_2 & \end{bmatrix}$
- 

The correctness of Algorithm 2 is stated in the following Theorem 2 and proven in the remaining of the current Section.

**THEOREM 2.** *Algorithm 2 correctly computes a symmetric indefinite PLDLTPT factorization revealing the rank profile matrix.*

**4.2.2 Second phase: off-diagonal pivoting.** Consider  $\mathbf{M} = \begin{bmatrix} \mathbf{0} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{C} \end{bmatrix}$ , where  $\mathbf{B} \in \mathbb{R}^{m \times n}$ , with  $m \leq n$ , has now an arbitrary rank  $r \leq m$ . Then its PLDUQ decomposition is of the form

$$\mathbf{B} = \mathbf{P} \begin{bmatrix} \mathbf{L}_1 \\ \mathbf{M}_1 \end{bmatrix} \begin{bmatrix} \mathbf{D}_1 \end{bmatrix} \begin{bmatrix} \mathbf{U}_1 & \mathbf{V}_1 \end{bmatrix} \mathbf{Q},$$

with  $\mathbf{D}_1$  diagonal, and  $\mathbf{L}_1$  and  $\mathbf{U}_1$  unit square triangular matrices, all three of order  $r$ . Then consider a conformal block decomposition of  $\mathbf{Q}\mathbf{C}\mathbf{Q}^T = \begin{bmatrix} \mathbf{C}_1 & \mathbf{C}_2 \\ \mathbf{C}_2^T & \mathbf{C}_3 \end{bmatrix}$  where  $\mathbf{C}_1$  is  $r \times r$ . It remains to eliminate  $\mathbf{C}_1$  and  $\mathbf{C}_2$  with the pivots found in  $\mathbf{B}$ , which leads to the following factorization:

$$\mathbf{M} = \begin{bmatrix} \mathbf{P} & \\ & \mathbf{Q}^T \end{bmatrix} \begin{bmatrix} \mathbf{L}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{M}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{G}_1 & \mathbf{U}_1^T & \mathbf{0} \\ \mathbf{G}_2 & \mathbf{V}_1^T & \mathbf{I} \end{bmatrix} \begin{bmatrix} \mathbf{0} & \mathbf{D}_1 & \mathbf{0} \\ \mathbf{D}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{Z} \end{bmatrix} \times \begin{bmatrix} \mathbf{L}_1^T & \mathbf{M}_1^T & \mathbf{G}_1^T & \mathbf{G}_2^T \\ \mathbf{0} & \mathbf{0} & \mathbf{U}_1 & \mathbf{V}_1 \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{I} \end{bmatrix} \begin{bmatrix} \mathbf{P}^T & \\ & \mathbf{Q} \end{bmatrix} \quad (1)$$

where  $\mathbf{G}_1$  satisfies

$$\mathbf{U}_1^T \mathbf{D}_1 \mathbf{G}_1^T + \mathbf{G}_1 \mathbf{D}_1 \mathbf{U}_1 = \mathbf{C}_1 \quad (2)$$

and  $\mathbf{G}_2 = (\mathbf{C}_2^T - \mathbf{V}_1^T \mathbf{D}_1 \mathbf{G}_1^T) \mathbf{U}_1^{-1} \mathbf{D}_1^{-1}$  and  $\mathbf{Z} = \mathbf{C}_3 - (\mathbf{V}_1^T \mathbf{D}_1 \mathbf{G}_2^T + \mathbf{G}_2 \mathbf{D}_1 \mathbf{V}_1)$ .

In order to produce an LDLT decomposition, there still remains to perform permutations to

- (1) compact the leading elements of the lower triangular matrix into a  $2r \times 2r$  invertible leading triangular submatrix,
- (2) make the  $\begin{bmatrix} \mathbf{D}_1 & \\ & \mathbf{D}_1 \end{bmatrix}$  matrix block diagonal with 1 or 2-dimensional diagonal blocks.

The permutation matrix

$$\mathbf{P}_c = \begin{bmatrix} \mathbf{I}_r & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{I}_{m-r} & \mathbf{0} \\ \mathbf{0} & \mathbf{I}_r & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{I}_{n-r} \end{bmatrix}, \quad (3)$$

corresponding to a block circular rotation, takes care of condition 1, while preserving precedence in the non-pivot rows. This is a requirement for the factorization to reveal the rank profile matrix [11, Theorem 20]. The decomposition becomes

$$\mathbf{N} = \begin{bmatrix} \mathbf{P} & \\ & \mathbf{Q}^T \end{bmatrix} \mathbf{P}_c \begin{bmatrix} \mathbf{L}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{G}_1 & \mathbf{U}_1^T & \mathbf{0} \\ \mathbf{M}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{G}_2 & \mathbf{V}_1^T & \mathbf{I} \end{bmatrix} \begin{bmatrix} \mathbf{0} & \mathbf{D}_1 & \mathbf{0} \\ \mathbf{D}_1 & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{Z} \end{bmatrix} \times \begin{bmatrix} \mathbf{L}_1^T & \mathbf{G}_1^T & \mathbf{M}_1^T & \mathbf{G}_2^T \\ \mathbf{0} & \mathbf{U}_1 & \mathbf{0} & \mathbf{V}_1 \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{I} \end{bmatrix} \mathbf{P}_c^T \begin{bmatrix} \mathbf{P}^T & \\ & \mathbf{Q} \end{bmatrix}. \quad (4)$$

In order to achieve Condition 2, we will transform the matrix  $\begin{bmatrix} \mathbf{0} & \mathbf{D}_1 \\ \mathbf{D}_1 & \mathbf{0} \end{bmatrix}$  into the block diagonal matrix  $\text{Diag} \left( \begin{bmatrix} \mathbf{0} & d_i \\ d_i & \mathbf{0} \end{bmatrix} \right)$  where  $d_i$  is the  $i$ th diagonal element in  $\mathbf{D}_1$ . To describe the process, we will focus on the matrix

$$\mathbf{M}_2 = \begin{bmatrix} \mathbf{L}_1 & \mathbf{0} \\ \mathbf{G}_1 & \mathbf{U}_1^T \end{bmatrix} \begin{bmatrix} \mathbf{0} & \mathbf{D}_1 \\ \mathbf{D}_1 & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{L}_1^T & \mathbf{G}_1^T \\ \mathbf{0} & \mathbf{U}_1 \end{bmatrix} = \bar{\mathbf{L}} \cdot \bar{\mathbf{D}} \cdot \bar{\mathbf{L}}^T,$$

and consider a splitting in halves of the matrix  $\mathbf{D}_1 = \begin{bmatrix} \mathbf{D}_{11} & \\ & \mathbf{D}_{12} \end{bmatrix}$  where  $\mathbf{D}_{11}$  has order  $r_1$  and  $\mathbf{D}_{12}$  order  $r_2$ . This leads to the conformal

decomposition

$$\begin{bmatrix} \mathbf{L}_{11} & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{L}_{12} & \mathbf{L}_{13} & \mathbf{0} & \mathbf{0} \\ \mathbf{G}_{11} & \mathbf{G}_{14} & \mathbf{U}_{11}^T & \mathbf{0} \\ \mathbf{G}_{12} & \mathbf{G}_{13} & \mathbf{U}_{12}^T & \mathbf{U}_{13}^T \end{bmatrix} \begin{bmatrix} \mathbf{0} & \mathbf{0} & \mathbf{D}_{11} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{D}_{12} \\ \mathbf{D}_{11} & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{D}_{12} & \mathbf{0} & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{L}_{11}^T & \mathbf{L}_{12}^T & \mathbf{G}_{11}^T & \mathbf{G}_{12}^T \\ \mathbf{0} & \mathbf{L}_{13}^T & \mathbf{G}_{14}^T & \mathbf{G}_{13}^T \\ \mathbf{0} & \mathbf{0} & \mathbf{U}_{11} & \mathbf{U}_{12} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{U}_{13} \end{bmatrix}$$

Then considering the permutation matrix

$$\mathbf{P}_d = \begin{bmatrix} \mathbf{I}_{r_1} & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{I}_{r_2} & \mathbf{0} \\ \mathbf{0} & \mathbf{I}_{r_1} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{I}_{r_2} \end{bmatrix},$$

one can form  $\mathbf{P}_d^T \bar{\mathbf{D}} \mathbf{P}_d = \begin{bmatrix} \mathbf{0} & \mathbf{D}_{11} & \mathbf{0} & \mathbf{0} \\ \mathbf{D}_{11} & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{D}_{12} \\ \mathbf{0} & \mathbf{0} & \mathbf{D}_{12} & \mathbf{0} \end{bmatrix}$  and  $\mathbf{P}_d^T \bar{\mathbf{L}} \mathbf{P}_d = \begin{bmatrix} \mathbf{L}_{11} & \mathbf{0} & \mathbf{0} & \mathbf{0} \\ \mathbf{G}_{11} & \mathbf{U}_{11}^T & \mathbf{G}_{14} & \mathbf{0} \\ \mathbf{L}_{12} & \mathbf{0} & \mathbf{L}_{13} & \mathbf{0} \\ \mathbf{G}_{12} & \mathbf{U}_{12}^T & \mathbf{G}_{13} & \mathbf{U}_{13}^T \end{bmatrix}$ .

Applying this process recursively changes  $\bar{\mathbf{D}}$  into the desired block diagonal form. Then the transformation of  $\bar{\mathbf{L}}$  will remain lower triangular if and only if all  $\mathbf{G}_{14}$  matrices are zero: this means that  $\mathbf{G}_1$  must be lower triangular in the first place.

Finding  $\mathbf{G}_1$  lower triangular satisfying Equation (2), is an instance of Problem 1 for which the routine `trssyr2k` provides a solution.

Note that the actual permutation to transform  $\begin{bmatrix} \mathbf{0} & \mathbf{D}_1 \\ \mathbf{D}_1 & \mathbf{0} \end{bmatrix}$  into a  $2 \times 2$ -blocks diagonal matrix is a permutation matrix,  $\mathbf{P}_i$ , resulting from the one by one interleaving of the rows of  $\begin{bmatrix} \mathbf{I}_r & \mathbf{0} \end{bmatrix}$  and  $\begin{bmatrix} \mathbf{0} & \mathbf{I}_r \end{bmatrix}$ . If  $\mathbf{e}_i = [0 \dots 0 \ 1 \ 0 \dots 0]^T$  is the  $i$ -th canonical vector, then:

$$\mathbf{P}_i = [\mathbf{e}_1 \ \mathbf{e}_{r+1} \ \mathbf{e}_2 \ \mathbf{e}_{r+2} \ \dots \ \mathbf{e}_r \ \mathbf{e}_{2r}]. \quad (5)$$

Similarly the triangular factor of the factorization is thus a one by one interleaving of the rows of  $\begin{bmatrix} \mathbf{L}_1 & \mathbf{0} \end{bmatrix}$  and  $\begin{bmatrix} \mathbf{G}_1 & \mathbf{U}_1^T \end{bmatrix}$  as well as a one by one interleaving of the columns  $\begin{bmatrix} \mathbf{L}_1 \\ \mathbf{G}_1 \end{bmatrix}$  and  $\begin{bmatrix} \mathbf{0} \\ \mathbf{U}_1^T \end{bmatrix}$ , which overall remains triangular.

Finally, a call to Algorithm 2 produces a factorization for the remaining  $\mathbf{Z}$  block and a final block rotation,

$$\begin{bmatrix} \mathbf{I}_{2r} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{I}_{m-r} \\ \mathbf{0} & \mathbf{I}_{n-r} & \mathbf{0} \end{bmatrix},$$

moves the intermediate zero rows and columns to the bottom right. The full algorithm is presented in details in Algorithm 3 (for zero or odd characteristic, the characteristic two case being presented afterwards in Section 4.3).

### 4.3 Characteristic two

The case of the characteristic two can be handled similarly, just computing the extra diagonal and updating after the PLDUQ decomposition, as sketched in Section 4.1.2. Indeed, the only issue is the division by 2 in `trssyr2k`, which is removed if the diagonal of  $\mathbf{C}'_1$  is zero. Therefore, Algorithm 2 is unchanged, the block diagonal matrix just has lower symmetric antitriangular  $2 \times 2$  blocks instead of only antidiagonal ones. The only few additional operations appear in Algorithm 3 (lines 4-9, 12, 16 and 21).

Then the tridiagonal form with symmetric antitriangular  $2 \times 2$  blocks thus obtained by Algorithm 3 can be used to either reveal the rank profile matrix (via computing  $\Psi_{\mathbf{D}}$ , the support matrix of  $\mathbf{D}$ , and the pivoting matrix  $\mathcal{R} = \mathbf{P} \Psi_{\mathbf{D}} \mathbf{P}^T$ ) or a PLDLTPT factorization, both at an extra linear cost in the dimension, as shown in Section 2.3.

**Algorithm 3** Rank deficient and zero leading principal symmetric elimination**Require:**  $C \in \mathbb{F}^{n \times n}$  symmetric and  $B \in \mathbb{F}^{m \times n}$ .**Ensure:**  $P$  permutation,  $L$  unit lower triangular,  $D$  block-diagonal, s.t.  $\begin{bmatrix} 0 & B \\ B^T & C \end{bmatrix} = PLDL^T P^T$ .

```

1: Decompose  $B = P_B \begin{bmatrix} L_1 \\ M_1 \end{bmatrix} \begin{bmatrix} D_1 & \\ & U_1 \end{bmatrix} Q$  {PLDUQ}
2: let  $r = \text{rank}(B)$  s.t.  $L_1, D_1$  and  $U_1$  are  $r \times r$ .
3:  $C' \leftarrow Q C Q^T = \begin{bmatrix} C'_1 & C'_2 \\ C'_2{}^T & C'_3 \end{bmatrix}$  where  $C'_1$  is  $r \times r$  {perm}
4: if  $\text{characteristic}(\mathbb{F}) = 2$  then
5:   for  $i = 1$  to  $r$  do
6:      $\Delta_{ii} = (C'_1)_{ii} - \sum_{j=1}^{i-1} \Delta_{jj} (U_1)_{ji}^2$ 
7:   end for
8:    $C'_1 \leftarrow C'_1 - U_1^T \Delta U_1$  {syrdk( $C'_1, U_1, \Delta$ )}
9: end if
10: Find  $X$  s.t.  $X^T U_1 + U_1^T X = C'_1$  {trssyr2k( $U_1, C'_1$ )}
11:  $G_1 \leftarrow X^T D_1^{-1}$  {scal( $X^T, D_1^{-1}$ )}
12: if  $\text{characteristic}(\mathbb{F}) = 2$  then  $X \leftarrow X + \Delta U_1$  end if {dadd( $X, \Delta, U_1$ )}
13:  $C'_2 \leftarrow C'_2 - X^T V_1$  {trmm( $X^T, V_1$ )}
14:  $Y \leftarrow U_1^{-T} C'_2$  {trsm( $U_1^T, C'_2$ )}
15:  $Z \leftarrow C'_3 - (Y^T V_1 + V_1^T Y)$  {syrdk( $C'_3, Y, V_1$ )}
16: if  $\text{characteristic}(\mathbb{F}) = 2$  then  $Z \leftarrow Z - (V_1^T \Delta V_1)$  end if {syrdk( $Z, V_1, \Delta$ )}
17:  $G_2 \leftarrow Y^T D_1^{-1}$  {scal( $Y^T, D_1^{-1}$ )}
18: Decompose  $Z = P_3 L_3 D_3 L_3^T P_3^T$  {Alg. 2}
19:  $P \leftarrow \begin{bmatrix} P_B & 0 \\ 0 & Q^T \end{bmatrix} \begin{bmatrix} P_c & 0 \\ 0 & I_{n-r} \end{bmatrix} \begin{bmatrix} P_i & 0 \\ 0 & I_{m+n-2r} \end{bmatrix} \begin{bmatrix} I_{m+r} & 0 \\ 0 & P_3 \end{bmatrix} \begin{bmatrix} I_{2r} & 0 & 0 \\ 0 & 0 & I_{m-r} \\ 0 & I_{n-r} & 0 \end{bmatrix}$  {With  $P_c$  from (3), and  $P_i$  from (5)}
20:  $L \leftarrow \begin{bmatrix} P_i^T \begin{bmatrix} L_1 \\ G_1 U_1^T \end{bmatrix} P_i & \\ \begin{bmatrix} G_2 V_1^T \\ M_1 & 0 \end{bmatrix} P_i & L_3 \\ & & 0 \end{bmatrix}$ 
21: if  $\text{characteristic}(\mathbb{F}) = 2$  then  $D \leftarrow \begin{bmatrix} P_i^T \begin{bmatrix} 0 & D_1 \\ D_1 & \Delta \end{bmatrix} P_i & \\ & D_3 \end{bmatrix}$  else  $D \leftarrow \begin{bmatrix} P_i^T \begin{bmatrix} 0 & D_1 \\ D_1 & 0 \end{bmatrix} P_i & \\ & D_3 \end{bmatrix}$  end if

```

**5 BASE CASE ITERATIVE VARIANT**

The recursion of Algorithm 2 should not be performed all the way to a dimension 1 in practice. For implementations over a finite field, it would induce an unnecessary large number of modular reductions and a significant amount of data movement for the permutations. Instead, we propose in Algorithm 4 an iterative algorithm computing a PLDLTPT revealing the rank profile matrix to be used as a base case in the recursion.

This iterative algorithm has the following features:

- (1) it uses a pivot search minimizing the lexicographic order (following the characterization in [11]): if the diagonal element of the current row is 0, the pivot is chosen as the first non-zero element of the row, unless the row is all zero, in which case, it is searched in the following row;
- (2) the pivot is permuted with cyclic shifts on the row and columns, so as to leave the precedence in the remaining rows and columns unchanged.
- (3) the update of the unprocessed part in the matrix is delayed following the scheme of a Crout elimination schedule [8]. It does not only improve efficiency thanks to a better data

locality, but it also reduces the amount of modular reductions, over a finite field, as shown for the unsymmetric case in [9].

We denote by  $\rho_{i,n}$  the cyclic shift permutation of order  $n$  moving element  $i$  to the first position:  $\rho_{i,n} = (i, 0, 1, \dots, i-1, i+1, \dots, n-1)$ . Indices are 0 based, index ranges are excluding their upper bound. For instance,  $A_{i,0..r}$  denotes the  $r$  first elements of the  $i+1$ st row of  $A$ , and  $A_{0..r,0..r}$  is the 0-dimensional matrix when  $r = 0$ .

**6 EXPERIMENTS**

We now report on experiments of an implementation of these algorithms in the FFLAS-FFPACK library [17], dedicated to dense linear algebra over finite fields. We used the version committed under the reference e12a998 of the master branch. It was compiled with gcc-5.4 and was linked with the numerical library OpenBLAS-0.2.18. Experiments are run on a single core of an Intel Haswell i5-4690, @3.5GHz. Computation speed is normalized as *effective Gfops*, an estimate of the number of field operations that an algorithm with classic matrix arithmetic would perform per second, divided by the computation time. For a matrix of order  $n$  and rank  $r$ , we defined this as:

$$\text{Effective Gfops} = (r^3/3 + n^2r - r^2n)/(10^9 \times \text{time}).$$

**Algorithm 4** SYTRF Crout iterative base case**Require:**  $A \in \mathbb{F}^{n \times n}$  symmetric**Ensure:**  $P$ , a permutation,  $L$ , unit lower triangular and  $D$ , block diagonal, such that  $A = PLDL^T P^T$ 

```

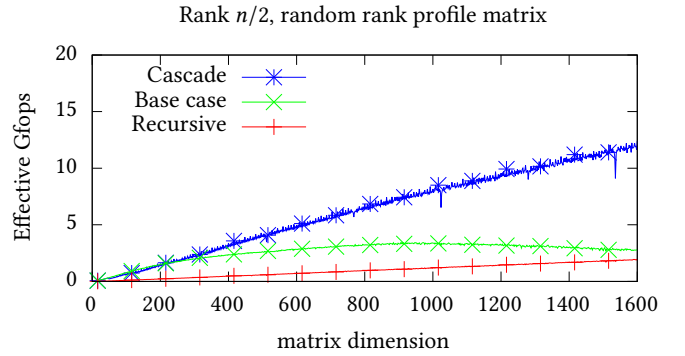
1: Denote  $W = A$  {the working matrix}
2:  $r \leftarrow 0$ ;  $D \leftarrow 0$ 
3: for  $i = 0..n$  do
4:   Here  $W = \begin{bmatrix} L & & \\ M & 0 & 0 \\ N & 0 & A_{i..n,i} & A_{i..n,i+1..n} \end{bmatrix}$  with  $L \in \mathbb{F}^{r \times r}$ 
5:    $v \leftarrow N_{0..r,i} \times D_{0..r,0..r}^{-1}$ 
6:    $c \leftarrow A_{i..n,i} - N \times v^T$ 
7:   if  $c = 0$  then Loop to next iteration end if
8:   Let  $j$  be the smallest index such that  $x = c_j \neq 0$ 
9:   Denote  $c = \begin{bmatrix} 0 & x & k \end{bmatrix}^T$ 
10:  if  $j = 0$  then
11:     $\begin{bmatrix} M \\ N \end{bmatrix} \leftarrow \rho_{j,n-r} \times \begin{bmatrix} M \\ N \end{bmatrix}$ 
12:     $W_{r..n,r} \leftarrow x^{-1} \times \rho_{j,n-r} \times \begin{bmatrix} 0 \\ c \end{bmatrix}$ 
13:     $P \leftarrow P \times \rho_{j,n-r}^T$ ;  $D_{r,r} \leftarrow x$ ;  $r \leftarrow r + 1$ 
14:  else
15:     $w \leftarrow N_{j,0..r} \times D_{0..r,0..r}^{-1}$ 
16:     $d \leftarrow A_{i..n,j+i} - N \times w^T (= \begin{bmatrix} 0 & x & g & y & h \end{bmatrix}^T)$ 
17:    Here  $W = \begin{bmatrix} L & & & & \\ M & 0 & 0 & 0 & 0 \\ & 0 & 0 & 0 & x & k^T \\ N & 0 & 0 & F & g & J^T \\ & 0 & x & g^T & y & h^T \\ & 0 & k & J & h & * \end{bmatrix}$ 
18:    if  $\text{characteristic}(\mathbb{F}) = 2$  then
19:       $y' \leftarrow 0$ ;  $h' \leftarrow h - yx^{-1}k$ 
20:       $D_{r..r+2,r..r+2} \leftarrow \begin{bmatrix} 0 & x \\ x & y \end{bmatrix}$ 
21:    else
22:       $y' \leftarrow y/2$ ;  $h' \leftarrow h - y'x^{-1}k$ 
23:       $D_{r..r+2,r..r+2} \leftarrow \begin{bmatrix} 0 & x \\ x & 0 \end{bmatrix}$ 
24:    end if
25:    Perform cyclic symmetric row and column rotations to
    bring  $W$  to the form  $W = \begin{bmatrix} L & & & & \\ n' & 1 & & & \\ & x^{-1}y' & 1 & & \\ M & 0 & 0 & 0 & 0 \\ N' & x^{-1}g & 0 & 0 & F & J^T \\ & x^{-1}h' & x^{-1}k & 0 & J & * \end{bmatrix}$ 
26:    Update  $P$  accordingly
27:     $r \leftarrow r + 2$ 
28:  end if
29: end for

```

All experiments are over the 23-bit finite field  $\mathbb{Z}/8388593\mathbb{Z}$  (timings with other primes of similar size are similar).

Figures 1 and 2 compare the computation speed of the pure recursive algorithm, the base case algorithm and a cascade of these two, with a threshold set to its optimum value from experiments

on this machine. Remark that the pure recursive variant performs rather well with generic rank profile matrices, while matrices with uniformly random rank profile matrix make this variant very slow, due to an excessive amount of pivoting. As expected, the base case Crout variant speeds up these instances for small dimensions, but then its performance stagnates on large dimensions, due to poor cache efficiency. Lastly the cascade algorithm combines the benefits of the two variants and therefore performs best in all settings. We here used a threshold  $n = 128$  for the experiments with random RPM matrices, but of only  $n = 48$  for generic rank profile matrices, since the recursive variant becomes competitive much earlier. In most cases, the rank profile structure of given matrices is unknown a priori, making the setting of this threshold speculative. One could instead implement an introspective strategy, updating the threshold from experimenting with running instances.



**Figure 1:** RPM revealing PLDLTPT: base case, pure recursive and cascading variants. Matrices with rank half the dimension and random RPM.

Table 1 compares the computation time of the finite field symmetric decomposition algorithm with several other routines: the unsymmetric elimination over a finite field (running the PLUQ algorithm of [11]) and the numerical elimination routines in double precision of LAPACK [2] provided by OpenBLAS: dgetrf (LU decomposition), and 3 variants of dsytrf [18] (symmetric LDLT decomposition with pivoting). These experiments first confirm a speed-up factor of about 2 between the symmetric and unsymmetric case over a finite field, which is the expected gain, looking at the constant in the time complexity. Note that on large instances, the PLUQ elimination performs better with random RPM instances than generic rank profiles, contrarily to the LDLT routine. This is due to the lesser amount of arithmetic operations when the RPM is random (some intermediate sub-matrices being rank deficient). On the other hand, these matrices generate more off-diagonal pivots, which cause more pivoting in LDLT than in PLUQ, explaining the slowdown for the symmetric case. On small dimensions, the numerical routines perform best as the quadratic modular reductions is penalizing the routines over a finite field. However, this is compensated by use of Strassen's algorithm, making our implementation outperform LAPACK's best dsytrf for larger dimension, even when the rank profile is not generic.




Figure 2: RPM revealing PLDLTPT. Matrices with full rank and random RPM (left) or generic rank profile (right).

$n$	Gen. rank prof. $r = n$				Gen. rank prof. $r = n$		Random RPM $r = n$		Random RPM $r = n/2$	
	dgetrf	dsytrf	dsytrf_rk	dsytrf_aa	PLUQ	LDLT	PLUQ	LDLT	PLUQ	LDLT
100	1.17e-04	1.31e-04	1.37e-04	9.21e-05	4.64e-04	3.23e-04	5.59e-04	5.22e-04	3.20e-04	3.80e-04
200	3.73e-04	4.39e-04	5.51e-04	4.48e-04	1.87e-03	8.80e-04	2.58e-03	1.59e-03	1.58e-03	1.33e-03
500	3.31e-03	3.78e-03	4.88e-03	3.87e-03	1.73e-02	5.21e-03	2.83e-02	9.85e-03	1.88e-02	7.92e-03
1000	2.19e-02	2.09e-02	2.58e-02	2.05e-02	8.84e-02	2.30e-02	9.95e-02	4.01e-02	6.27e-02	3.18e-02
2000	0.145	0.127	0.154	0.127	0.438	0.127	0.490	0.191	0.274	0.150
5000	2.005	1.604	1.871	1.598	3.904	1.591	3.849	1.744	2.431	1.294
10000	14.948	11.981	13.396	12.008	24.115	10.904	23.985	11.209	14.775	7.894

Table 1: Comparing computation time (s) of numerical routines with the symmetric (LDLT) and unsymmetric (PLUQ) triangular decompositions. Matrices with rank  $r$ , generic rank profile or rank profile matrix uniformly random.

## ACKNOWLEDGMENT

We thank the referees for their valuable feedback and pointing out the recent work of [18].

## REFERENCES

- [1] J. O. Aasen. On the reduction of a symmetric matrix to tridiagonal form. *BIT Numerical Mathematics*, 11(3):233–242, Sep 1971. doi:10.1007/BF01931804.
- [2] E. Anderson, Z. Bai, C. Bischof, L. S. Blackford, J. Demmel, J. Dongarra, J. Du Croz, A. Greenbaum, S. Hammarling, A. McKenney, et al. *LAPACK Users' guide*. SIAM, 1999. URL: [http://www.netlib.org/lapack/lug/lapack\\_lug.html](http://www.netlib.org/lapack/lug/lapack_lug.html).
- [3] M. Baboulin, D. Becker, and J. Dongarra. A Parallel Tiled Solver for Dense Symmetric Indefinite Systems on Multicore Architectures. In *IEEE 26th International Parallel & Distributed Processing Symposium (IPDPS)*, pages 14–24. IEEE, May 2012. doi:10.1109/IPDPS.2012.12.
- [4] G. Ballard, D. Becker, J. Demmel, J. Dongarra, A. Druinsky, I. Peled, O. Schwartz, S. Toledo, and I. Yamazaki. Communication-Avoiding Symmetric-Indefinite Factorization. *SIAM Journal on Matrix Analysis and Applications*, 35(4):1364–1406, Jan. 2014. doi:10.1137/130929060.
- [5] J. R. Bunch and L. Kaufman. Some stable methods for calculating inertia and solving symmetric linear systems. *Mathematics of Computation*, 31(137):163–179, 1977. doi:10.2307/2005787.
- [6] J. R. Bunch and B. N. Parlett. Direct methods for solving symmetric indefinite systems of linear equations. *SIAM Journal on Numerical Analysis*, 8(4):639–655, Dec. 1971. doi:10.1137/0708060.
- [7] J. J. Dongarra, J. Du Croz, S. Hammarling, and I. S. Duff. A Set of Level 3 Basic Linear Algebra Subprograms. *ACM TOMS*, 16(1):1–17, Mar. 1990. doi:10.1145/77626.79170.
- [8] J. J. Dongarra, L. S. Duff, D. C. Sorensen, and H. A. V. Vorst. *Numerical Linear Algebra for High Performance Computers*. SIAM, 1998.
- [9] J.-G. Dumas, T. Gautier, C. Pernet, J.-L. Roch, and Z. Sultan. Recursion based parallelization of exact dense linear algebra routines for gaussian elimination. *Parallel Computing*, 57:235–249, 2016. doi:10.1016/j.parco.2015.10.003.
- [10] J.-G. Dumas, C. Pernet, and Z. Sultan. Computing the rank profile matrix. In *Proceedings of the 2015 ACM on International Symposium on Symbolic and Algebraic Computation, ISSAC '15*, pages 149–156, New York, NY, USA, 2015. ACM. doi:10.1145/2755996.2756682.
- [11] J.-G. Dumas, C. Pernet, and Z. Sultan. Fast computation of the rank profile matrix and the generalized Bruhat decomposition. *Journal of Symbolic Computation*, 83:187–210, Nov.–Dec. 2017. doi:10.1016/j.jsc.2016.11.011.
- [12] E. Elmroth, F. G. Gustavson, I. Jonsson, and B. Kågström. Recursive blocked algorithms and hybrid data structures for dense matrix library software. *SIAM Review*, 46(1):3–45, 2004. doi:10.1137/S0036144503428693.
- [13] E. L. Kaltofen, M. Nehring, and B. D. Saunders. Quadratic-time certificates in linear algebra. In A. Leykin, editor, *ISSAC'2011, Proceedings of the 2011 ACM International Symposium on Symbolic and Algebraic Computation, San Jose, California, USA*, pages 171–176. ACM Press, New York, June 2011. doi:10.1145/1993886.1993915.
- [14] B. Parlett and J. K. Reid. On the solution of a system of linear equations whose matrix is symmetric but not definite. *BIT*, 10(3):386–397, 1970. doi:10.1007/BF01934207.
- [15] M. Rozložník, G. Shklarski, and S. Toledo. Partitioned triangular tridiagonalization. *ACM Trans. Math. Softw.*, 37(4):38:1–38:16, Feb. 2011. doi:10.1145/1916461.1916462.
- [16] G. Shklarski and S. Toledo. Blocked and recursive algorithms for triangular tridiagonalization. 2007. URL: <http://www.cs.tau.ac.il/~stoledo/Bib/Pubs/ShklarskiToledo-Aasen.pdf>.
- [17] The FFLAS-FFPACK group. *FFLAS-FFPACK: Finite Field Linear Algebra Subroutines / Package*, 2018. v2.3.2. <https://github.com/linbox-team/fflas-ffpack>.
- [18] I. Yamazaki and J. Dongarra. Aasen's symmetric indefinite linear solvers in LAPACK. Technical Report 294, LAPACK Working Note, Dec. 2017. URL: <http://www.netlib.org/lapack/lawnspdf/lawn294.pdf>.